

ASP.NET CORE MICROSERVICES



BY:

MAHESH SABNIS

[HTTPS://WWW.WEBNETHELPER.COM](https://www.webnethelper.com)

[WWW.DEVCURRY.COM](http://www.devcurry.com)

[WWW.DOTNETCURRY.COM](http://www.dotnetcurry.com)



WHAT ARE MICROSERVICES?

The design requirement from the software system




WHAT ARE MICROSERVICES?

....and what we design and deliver



WHAT ARE MICROSERVICES?

How each tool looks like from inside



WHAT ARE MICROSERVICES

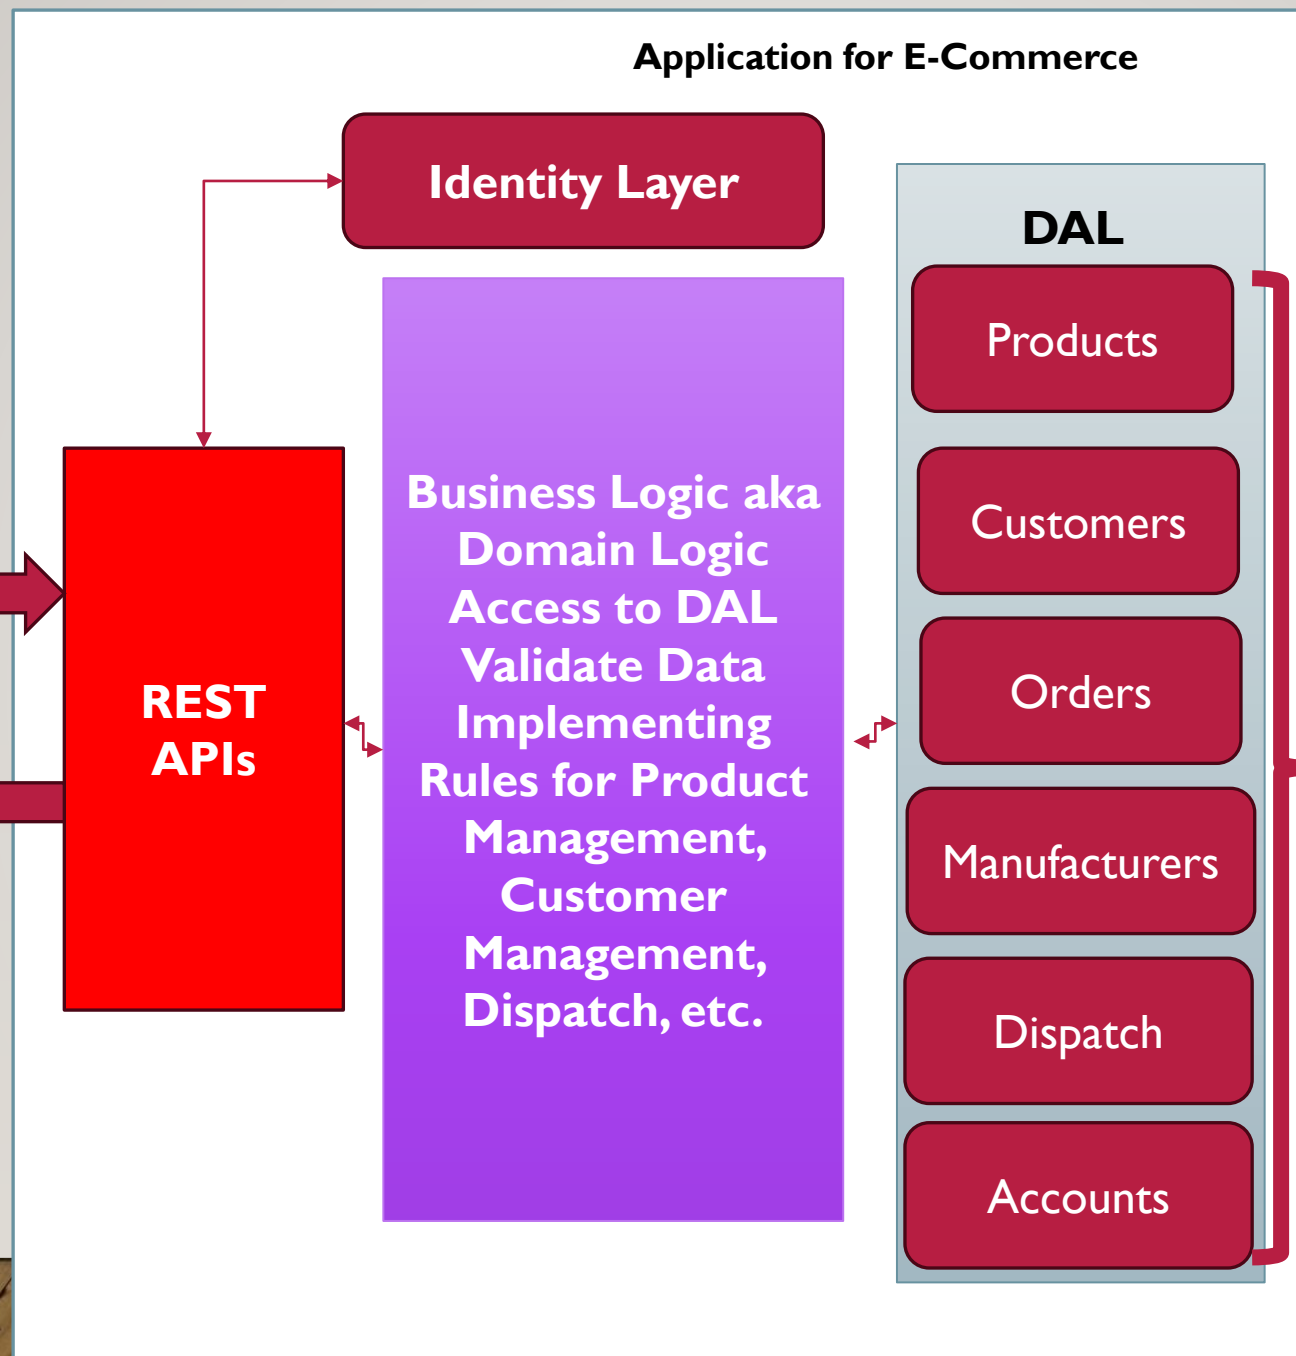
What is Microservice?

Isn't this just another service-oriented architecture (SOA) or domain-driven design (DDD) approach?


You can think of microservices as "SOA done right," with principles and patterns like autonomous services, Bounded-Context pattern and event-driven all having their roots in SOA and DDD.

WHAT ARE MICROSERVICES?

- The microservice architectural style is an approach to developing a single application as a **suite of small services**, each **running in its own process** and communicating with lightweight mechanisms, often an HTTP resource API. These services are **built around business capabilities** and **independently deployable** by fully automated deployment machinery. There is a **bare minimum of centralized management** of these services, which may be written in different programming languages and use different data storage technologies.
- -- [James Lewis and Martin Fowler \(2014\)](#)

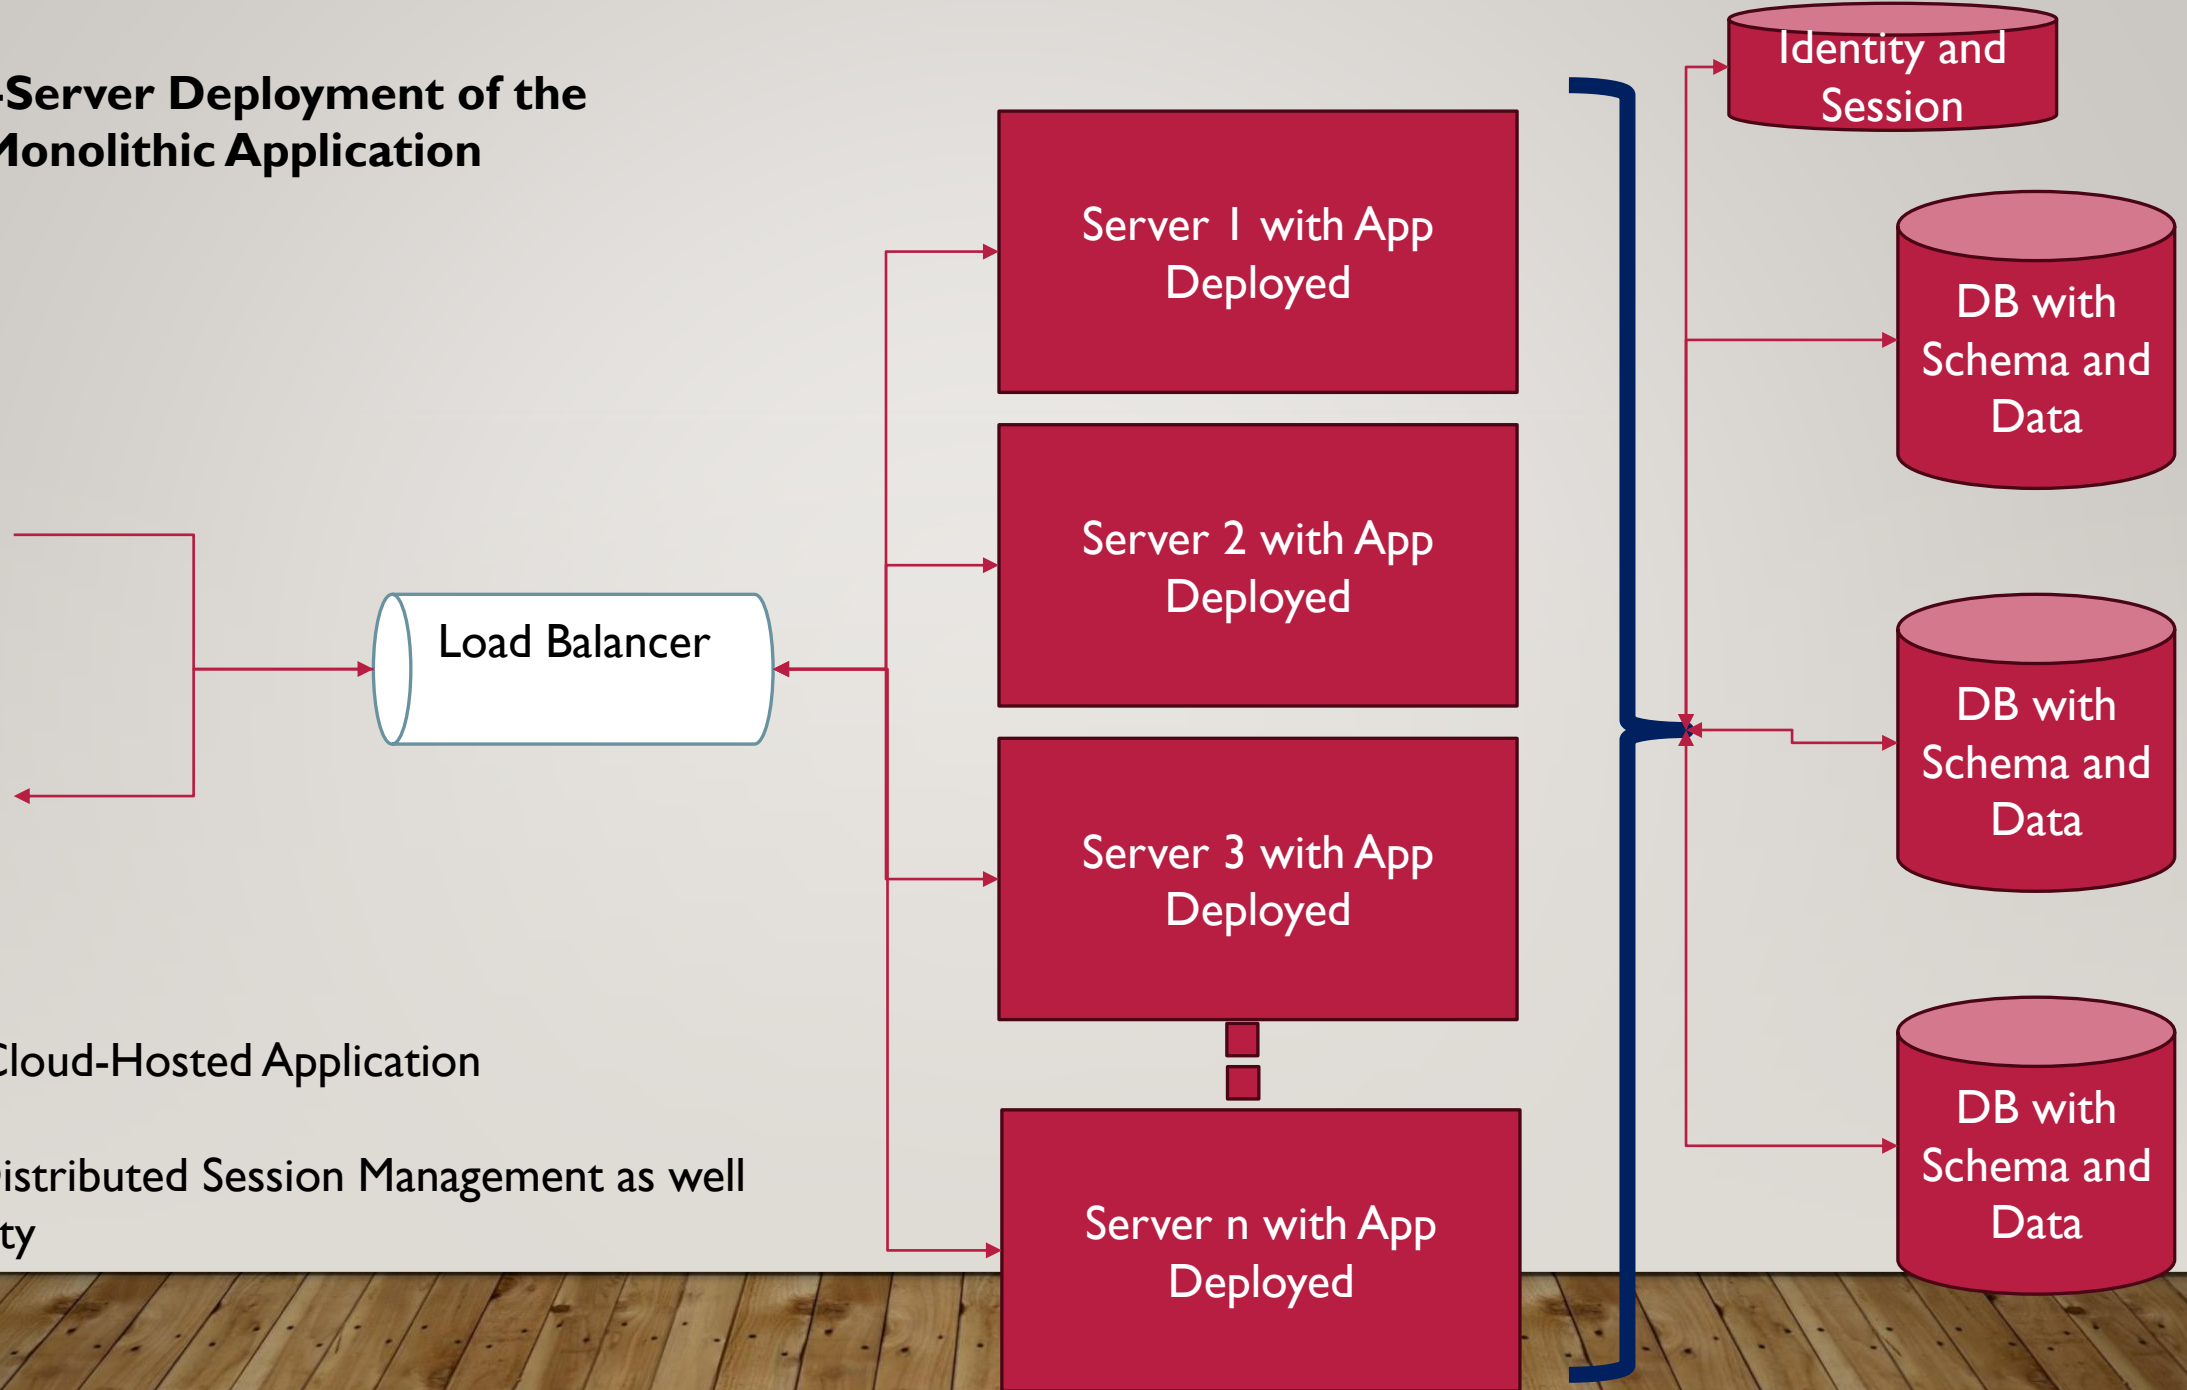


E-Commerce Database
with the various tables
as per the domain



**Monolithic
Applications**


Multi-Server Deployment of the Monolithic Application



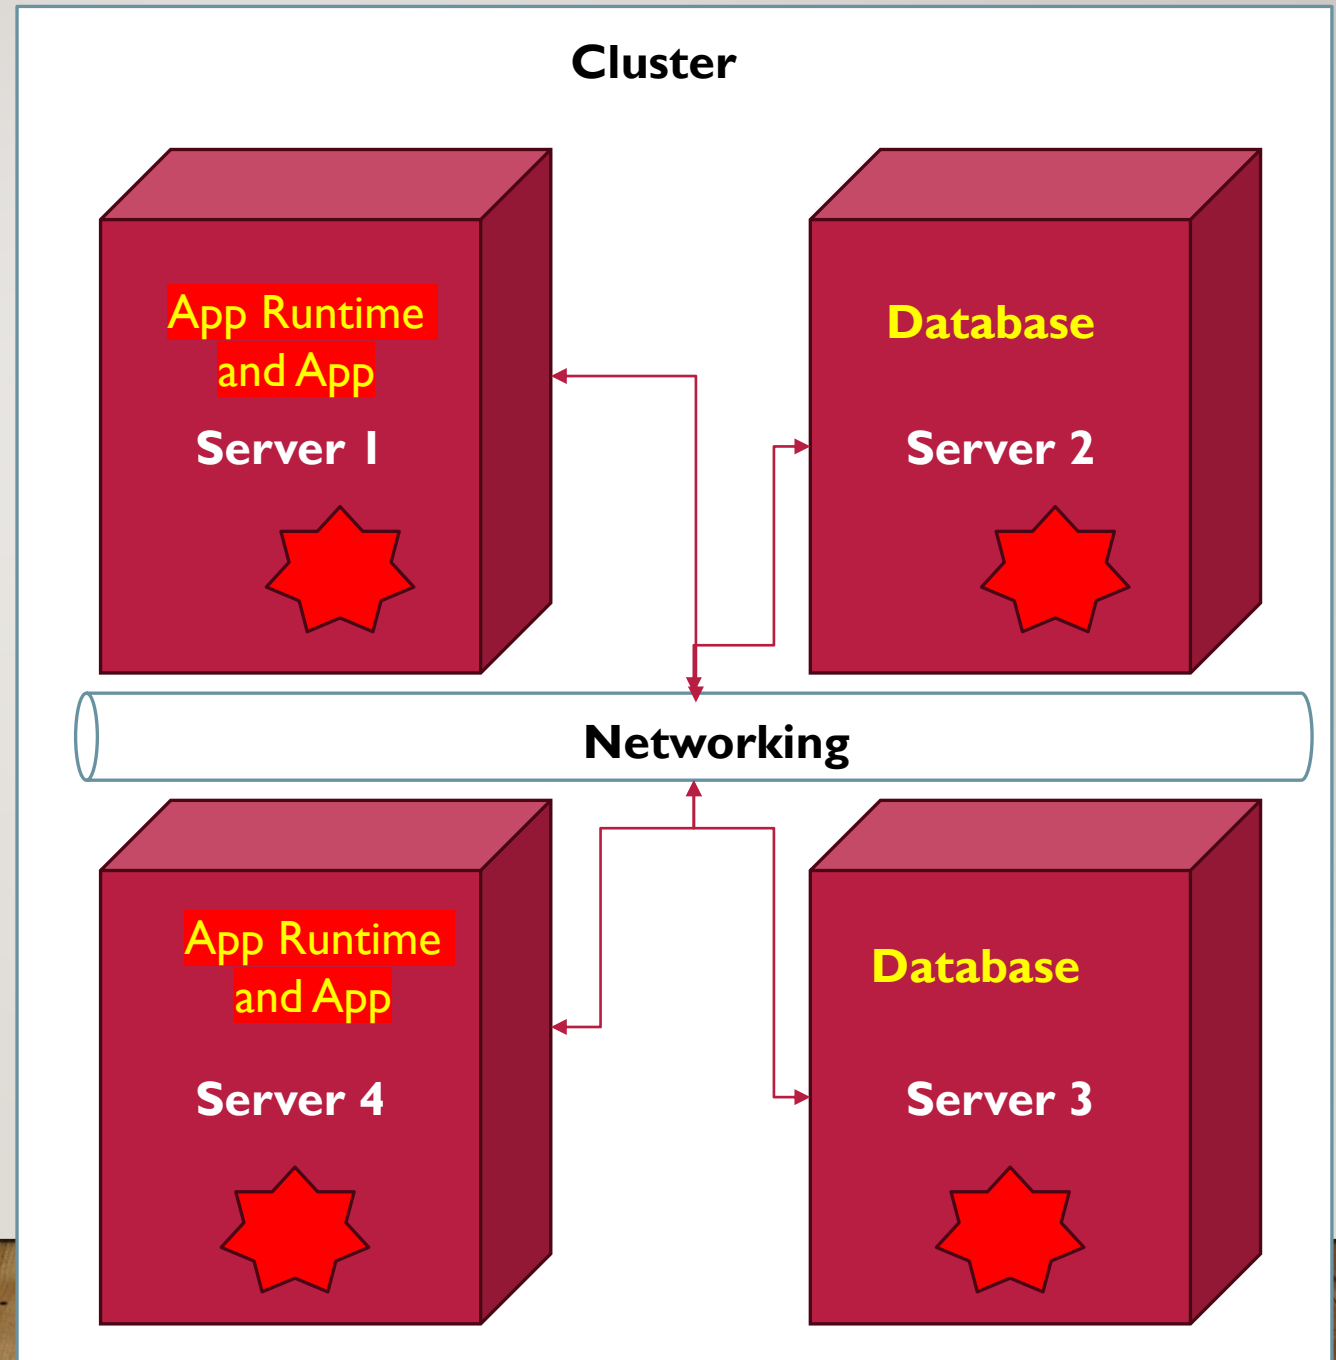
Private Cloud-Hosted Application

Needs Distributed Session Management as well as Security

**Separate
Each
Workflow as
independent
Application**



1. Multi-Server Deployment
 1. Each Server MUST be configured separately
 2. Each server MUST have an App deployed on it
 3. App Dependencies MUST be configured on each Server
2. State of the user (if needed) MSUT be saved in distributed storage

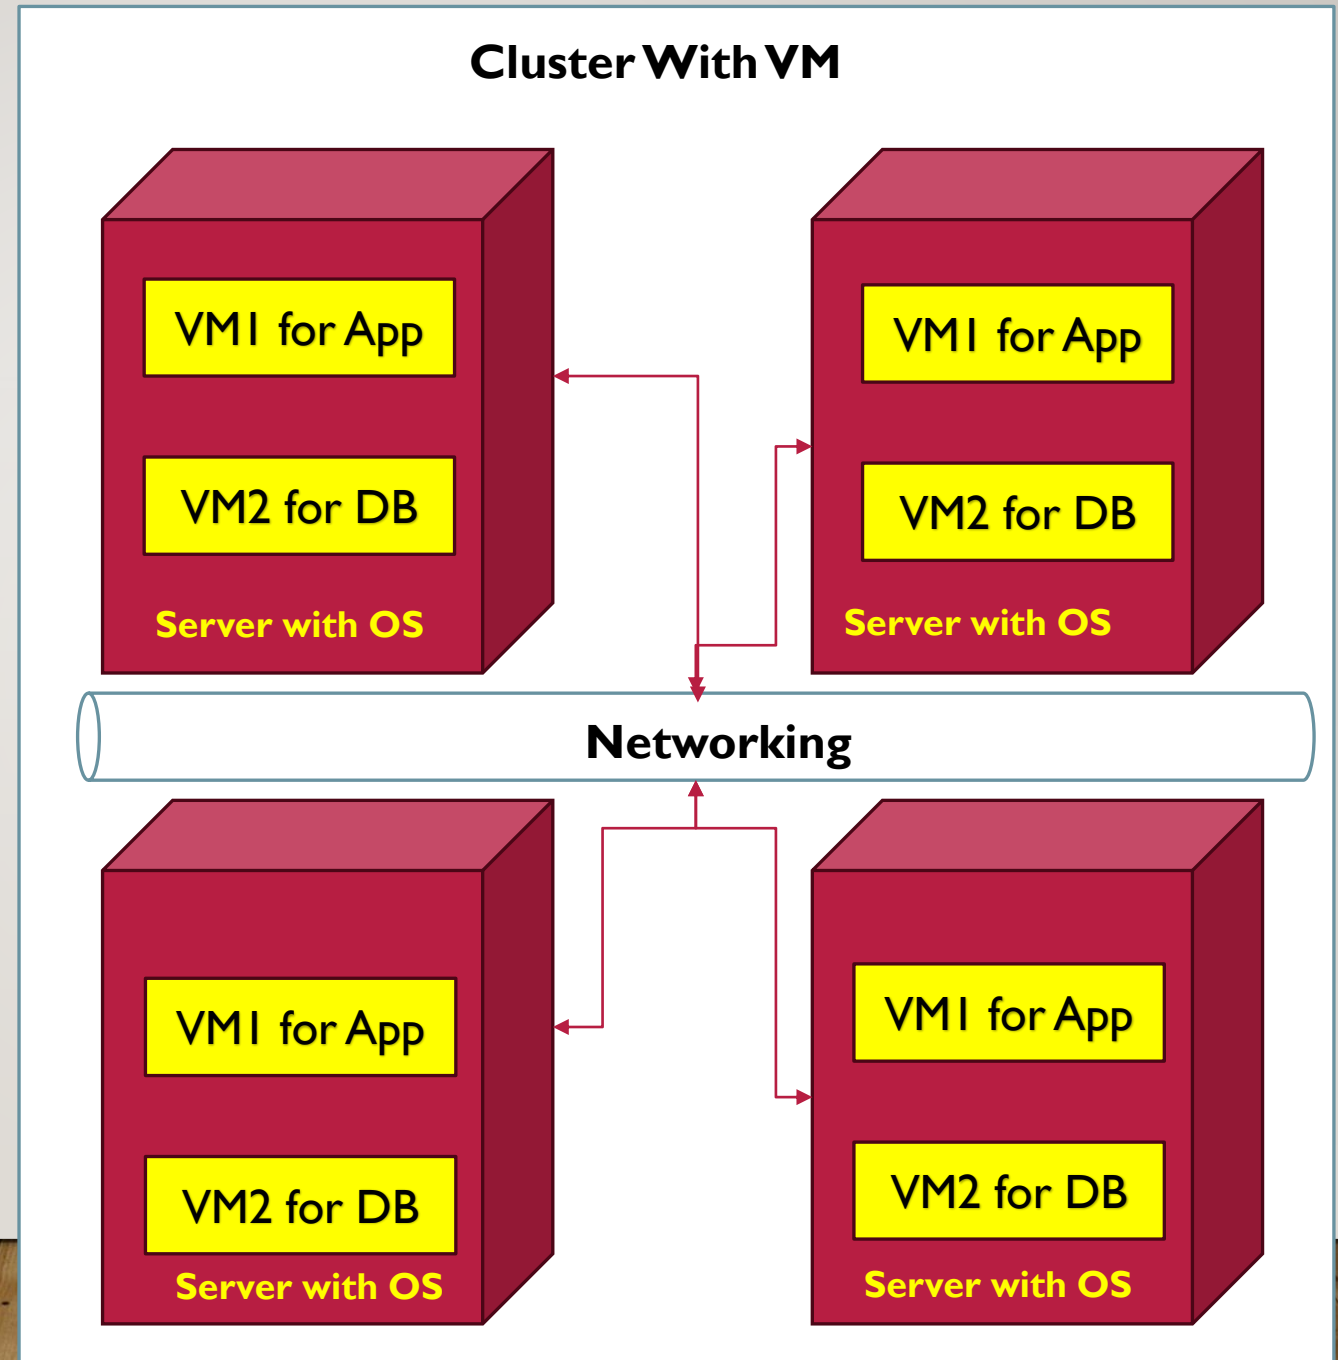


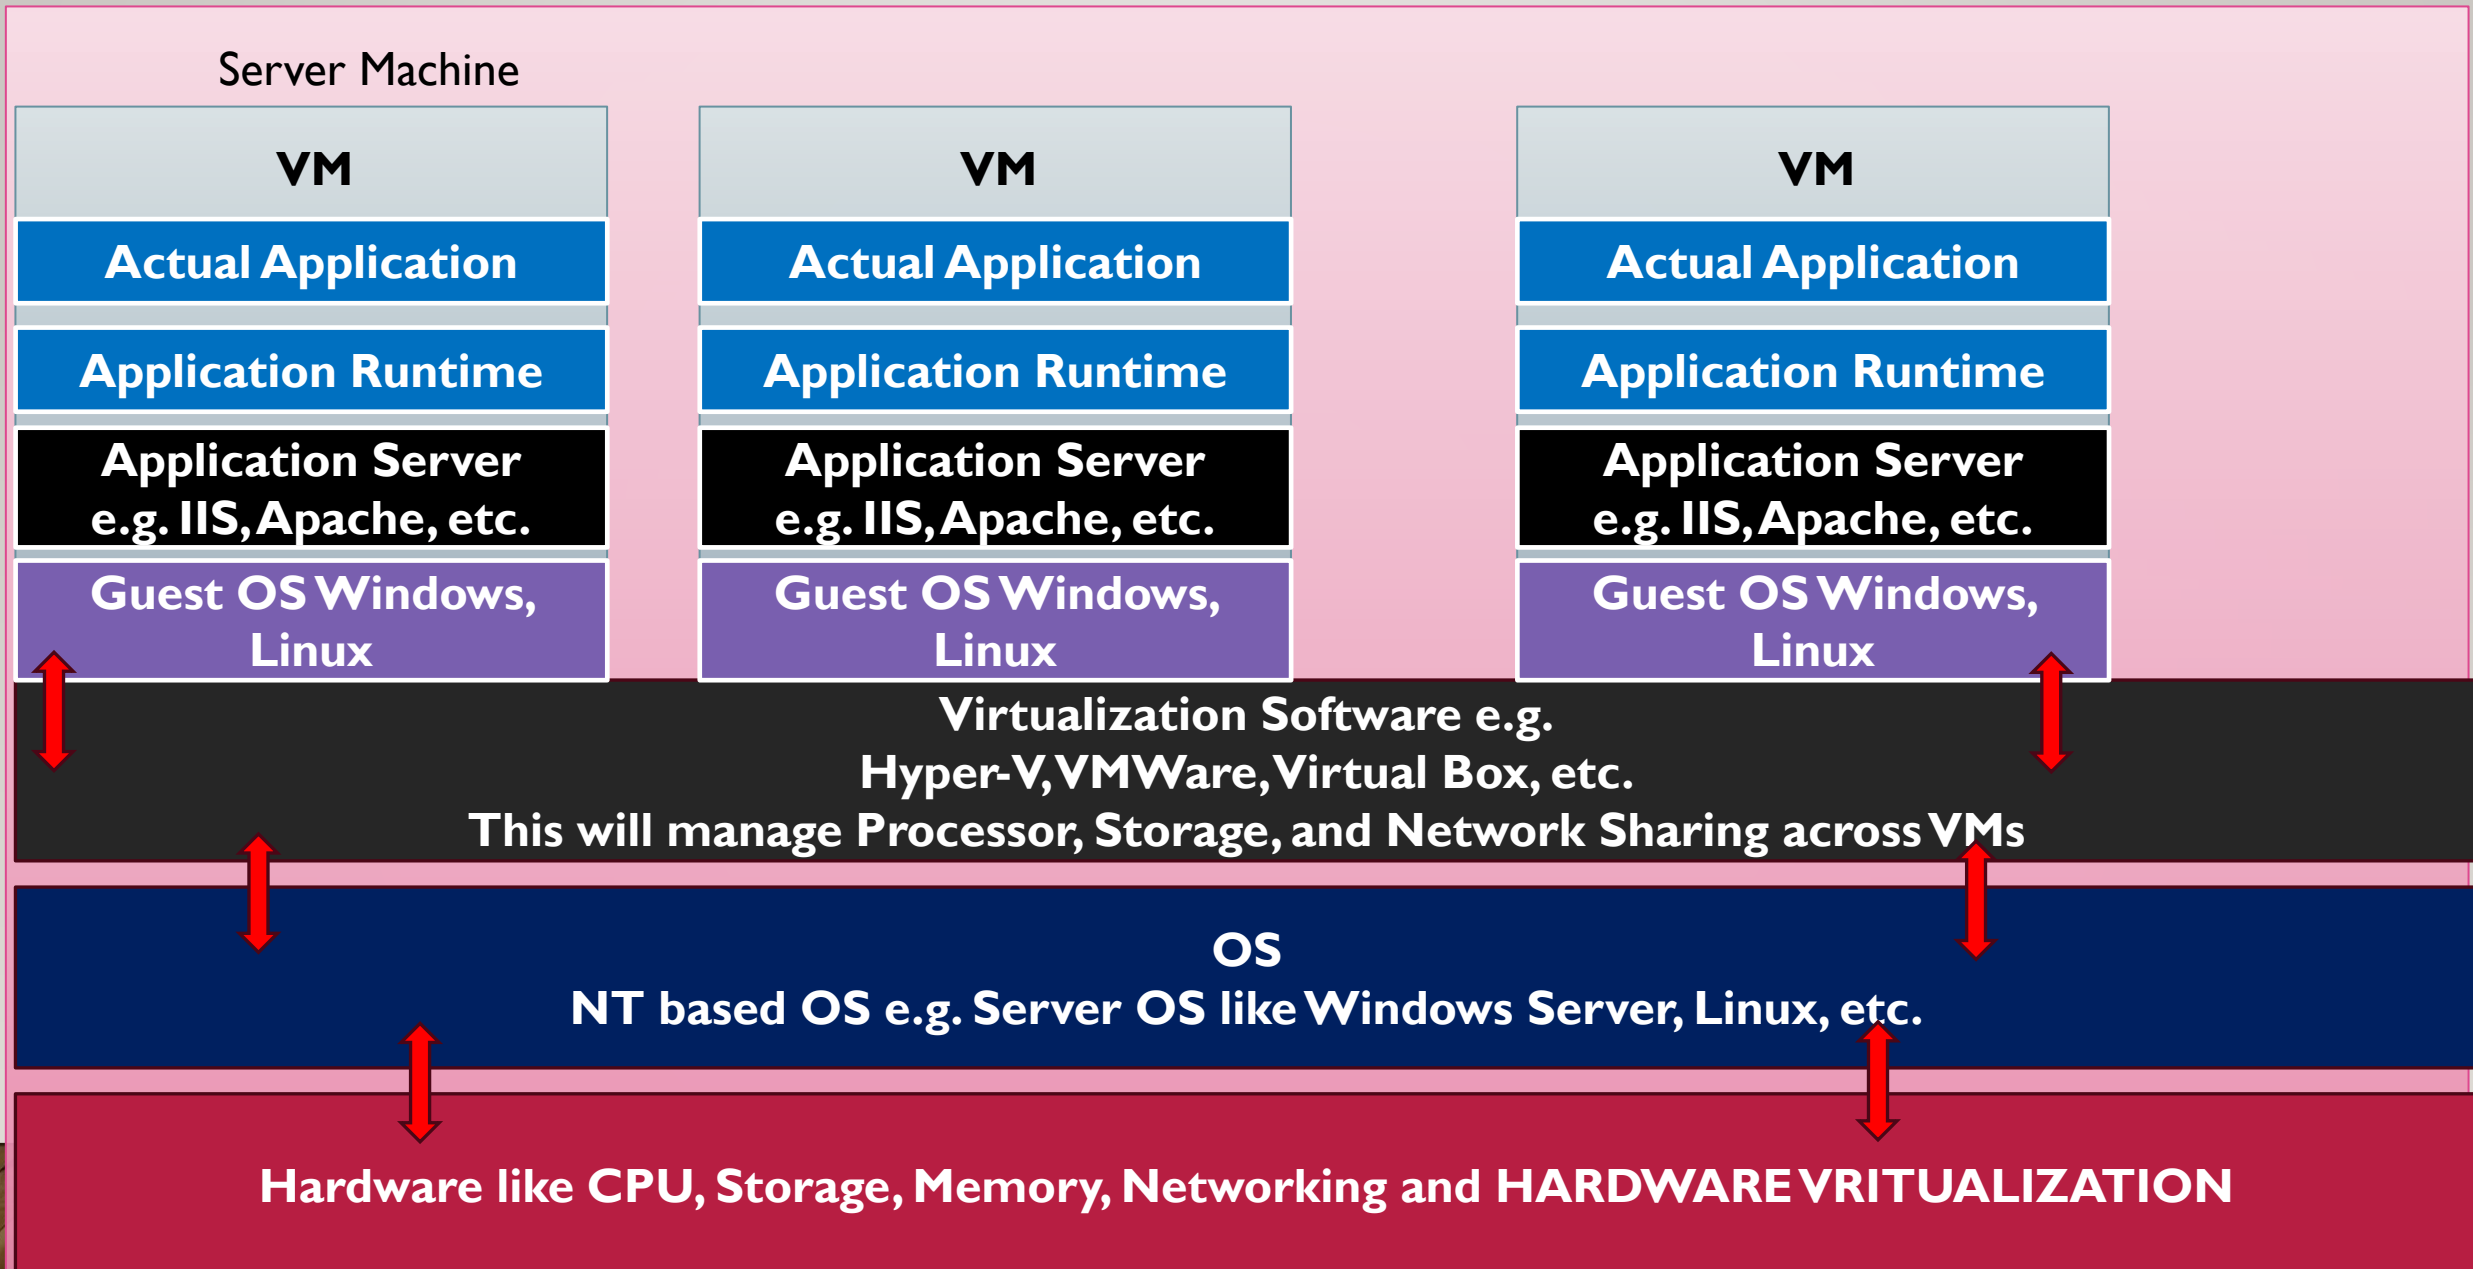
Generally, this is used by Cloud Enables Applications

Each Server MUST have Host OS and It MUST have Virtualization Software to run VM

VM MUST have

- Guest OS
- Application Runtime (If VM for App)
 - Actual App
- Database (if VM for DB)





The Containerization on the Physical machine (Windows, Linux, macOS)

Docker Container

Image of the Actual Application with all of its dependencies e.g. Runtime, and other third party dependencies

Image for OS to Run App

Docker Container

Image of the Actual Application with all of its dependencies e.g. Runtime, and other third party dependencies

Image for OS to Run App

Docker Container

Image of the Actual Application with all of its dependencies e.g. Runtime, and other third party dependencies

Image for OS to Run App

Shared Resources e.g. CPU, Memory, Storage, and Networking

The Container Hosting Service Service named as DOCKER
For Windows we need of Hyper-V

OS

NT based OS e.g. Server OS like Windows Server, Linux, etc.

Hardware like CPU, Storage, Memory, Networking and HARDWARE VIRTUALIZATION

WHAT ARE MICROSERVICES?

.....SOME TECHNICAL SPECIFICATIONS

Autonomous

Isolated

Elastic

Resilient

Responsive

Automated

Programmable

Message
Oriented

Intelligent

Configurable

Discoverable

Domain driven

WHAT ARE MICROSERVICES?

.....SOME BENEFITS

Evolutionary

Owned

Small

Safe

Versioned

Replaceable

Scale
Governance

Deployment
Governance

Reuse

Speed of
Development

Resilient

Open

WHAT ARE MICROSERVICES?

.....SOME CHALLENGES

Evolutionary

Discoverability

Testing

Domain
Modeling

Versioning
must be
Supported

Platform
Matters

Automation is
not an Option

Communication
is Key

MICROSERVICES DESIGN SOME THOUGHTS



THINKING FOR DOMAIN MODELING IN EXISTING SYSTEM

Bounded Context	<ul style="list-style-type: none">• Find the capability within the system - Where does the language change?
Rate of Change	<ul style="list-style-type: none">• How fast or slow areas of the system change can identify a seam
Team Structure	<ul style="list-style-type: none">• Service boundaries defined by how the org is structured
Pain	<ul style="list-style-type: none">• Whatever hurts most is a good candidate
Namespace Modeling	<ul style="list-style-type: none">• Partition by existing coupling



GOING FOR NEW DESIGN -- DOMAIN MODELING – GREEN FIELD

User Stories

- Leverage User Journey to find the capabilities

Evolutionary

- One service at a time

Single Responsibility

- High Cohesion

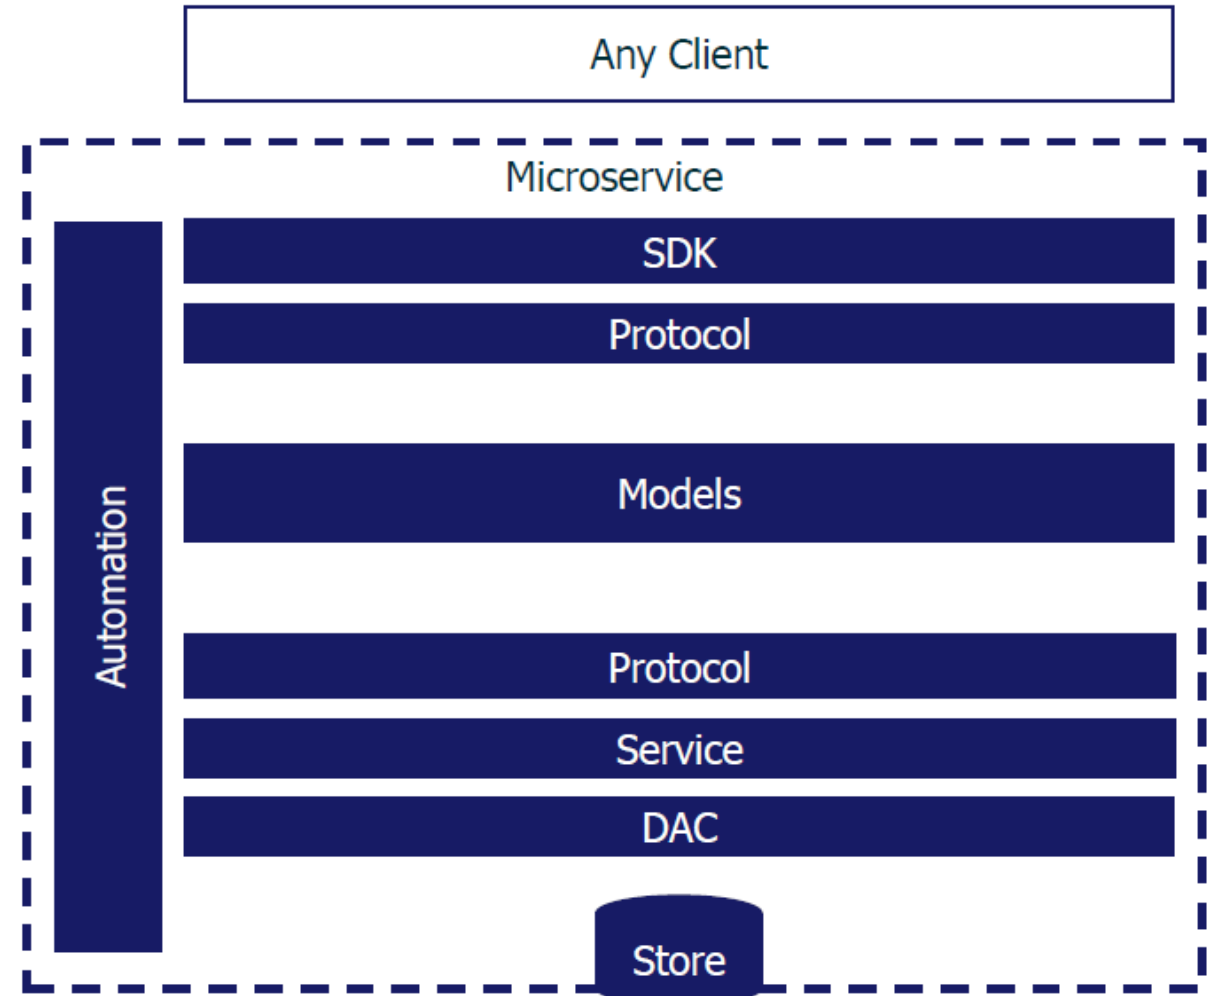
Loose Coupling

- Minimize dependencies

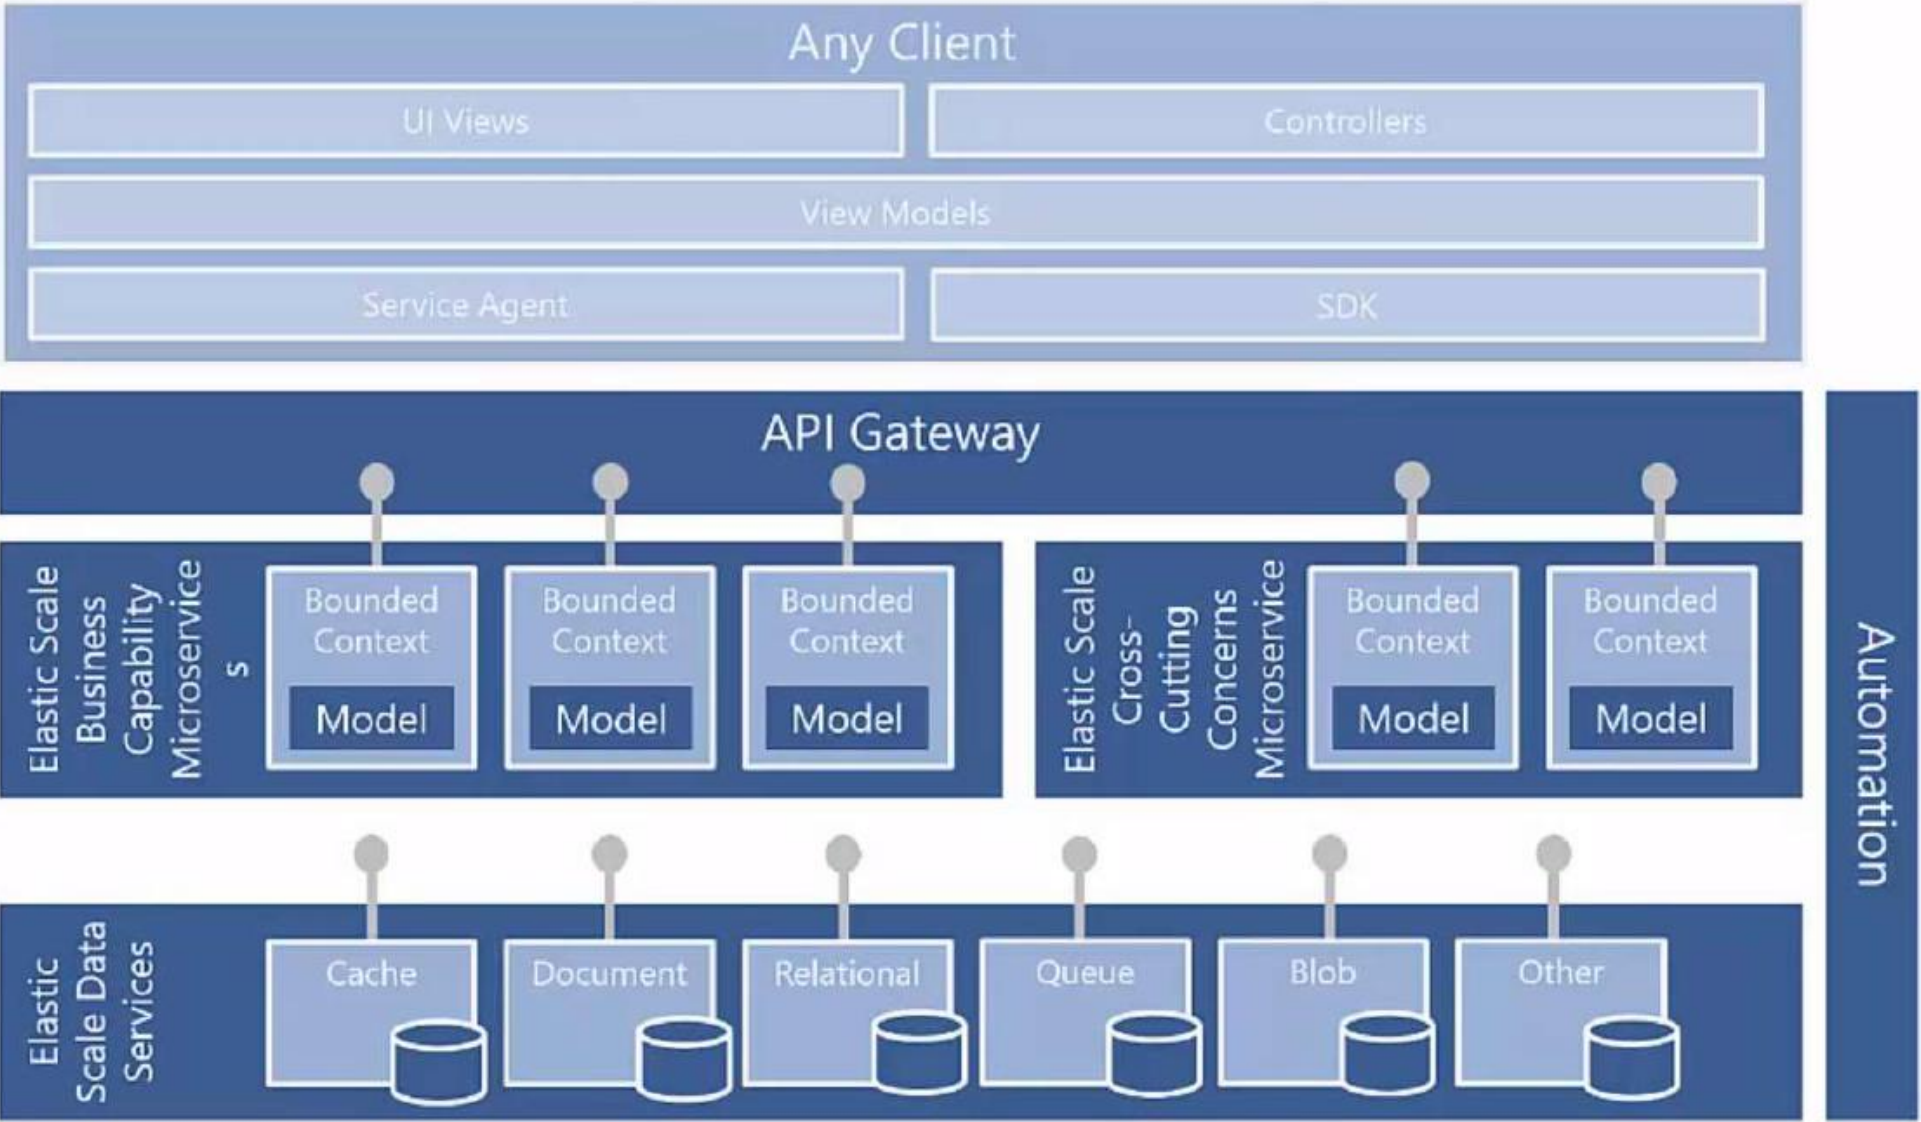
MICROSERVICES LOGICAL ARCHITECTURE FOR DOMAIN DRIVEN DEVELOPMENT



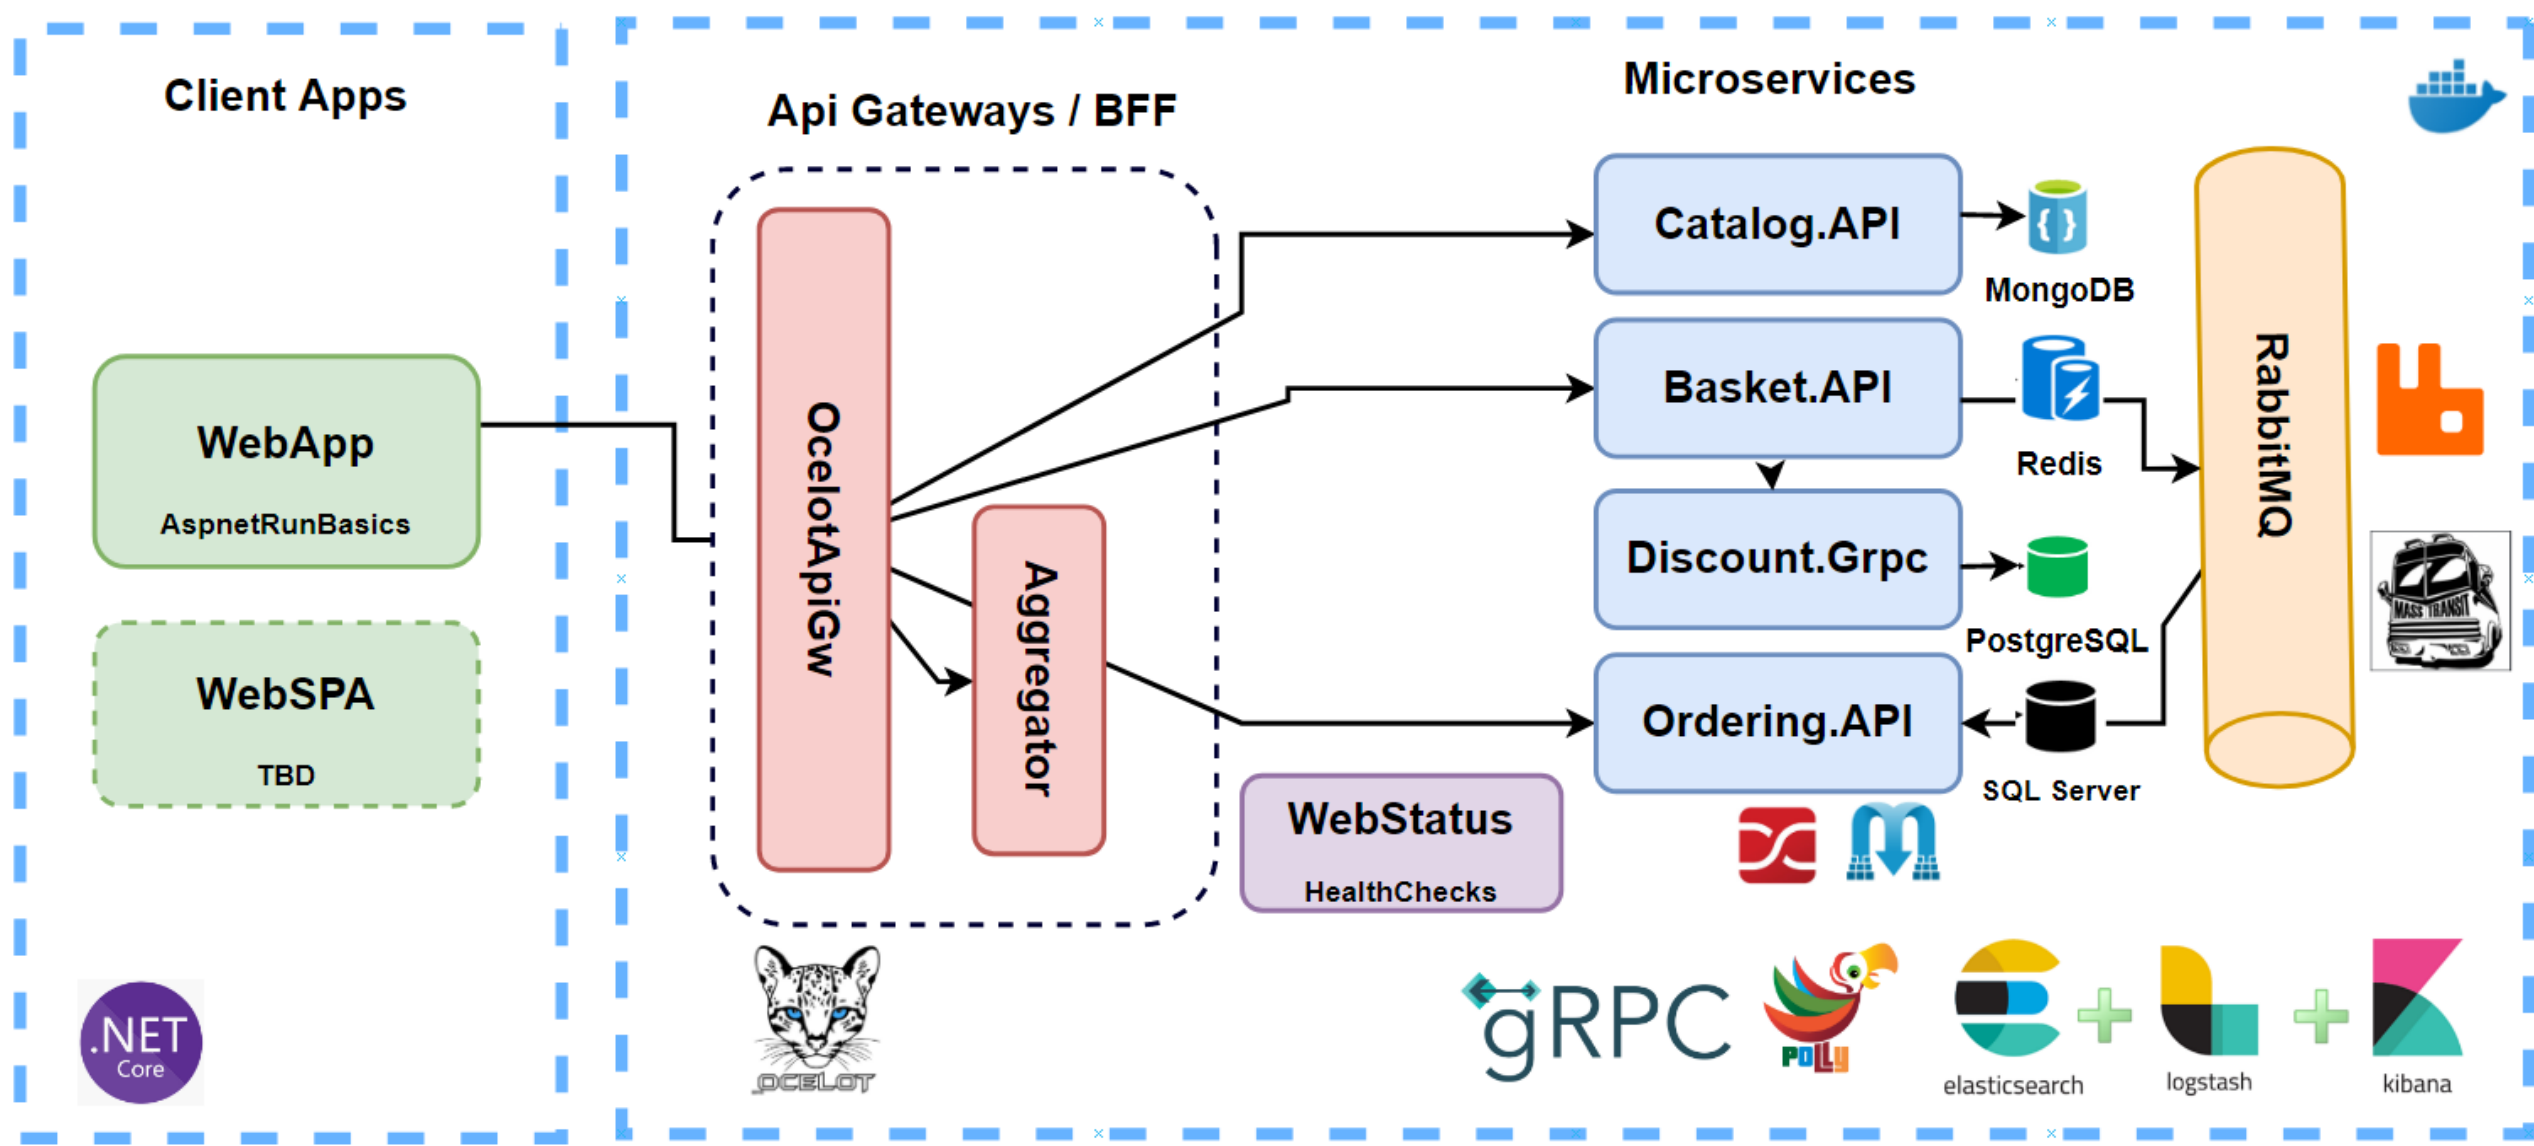
LOGICAL ARCHITECTURE



SERVER-SIDE DISCOVERY PATTERN



ASP.NET CORE MICROSERVICES



THE 'STATE' IN MICROSERVICE

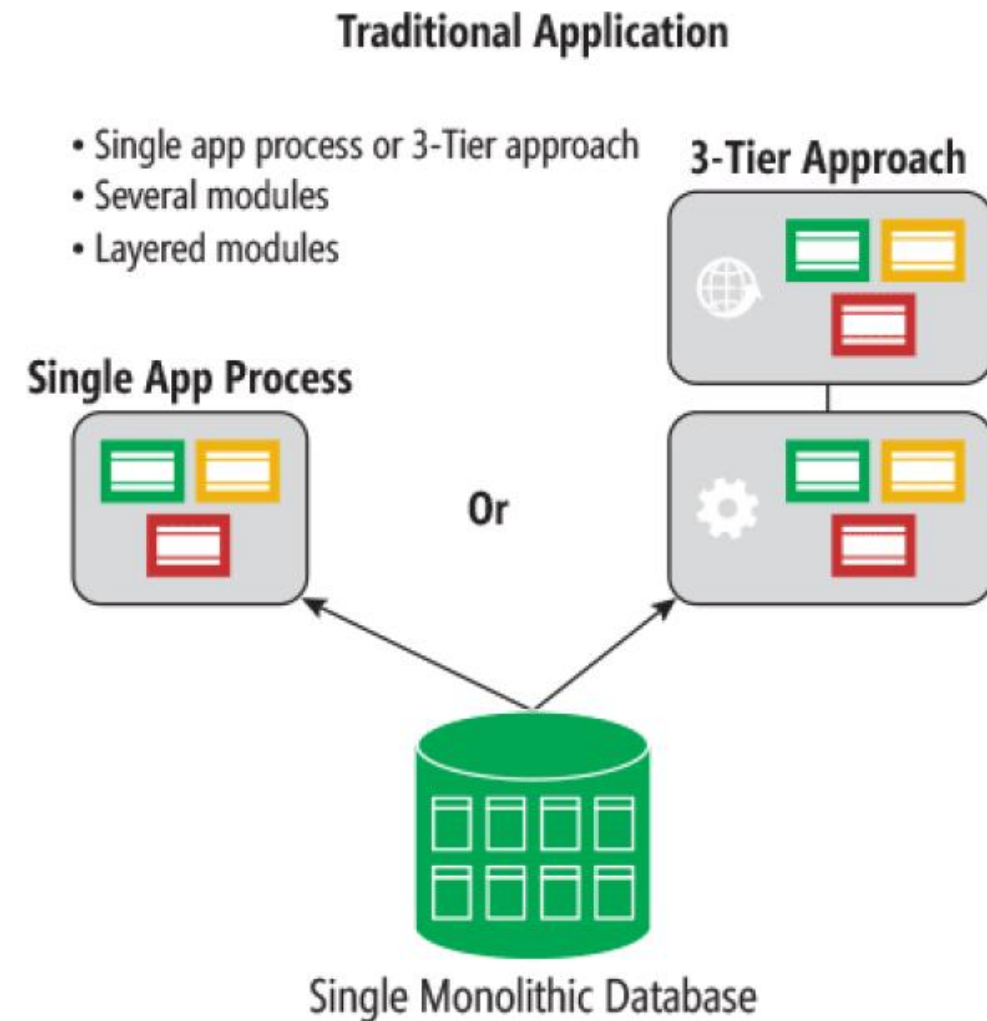
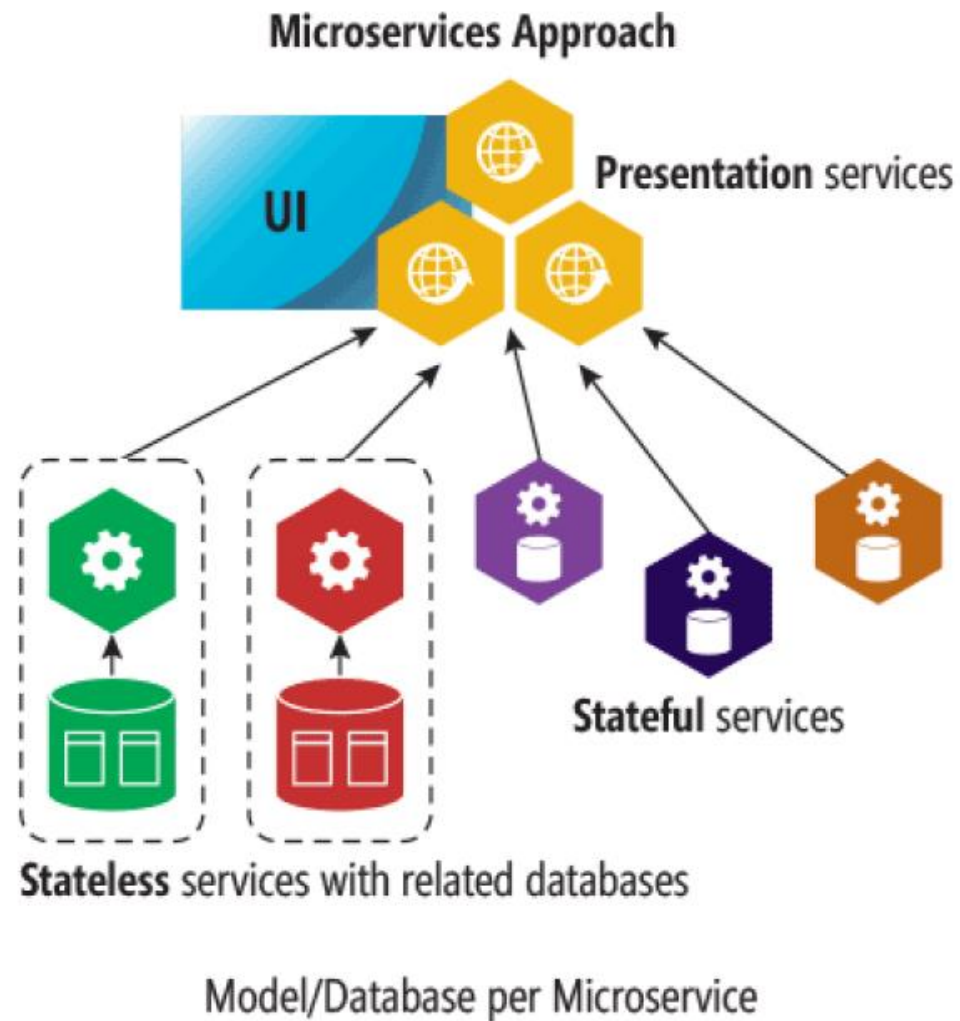
- What is State means in case of the Microservices?
- The data
 - Send to Microservices
 - Data processed and Generated by Microservices during processing
- Case to maintain the state
 - Microservices needs to maintain the data for high data availability
 - Data processing performance improvements
 - Data Consistency
 - Transaction Management
 - Data Reliability



THE 'STATE' IN MICROSERVICE

- Stateless Microservices
 - They need an external data repository to maintain the state of the data
 - E.g. Cache, Message Queues, File System, Database etc.
 - The data is stored out of the process of the microservices
 - Need to arrange and configure external store explicitly
 - Challenges in reliability because of the external dependency
- Stateful Microservices
 - Data is maintained in the local process of the Microservice
 - The cluster manager is responsible to transfer the data state
 - Reliability and availability is controlled by the cluster manager
 - Cluster manager comes with additional cost

STATEFUL VS STATELESS VS MONOLITHIC APPLICATIONS



PROGRAMMING WITH MICROSERVICES



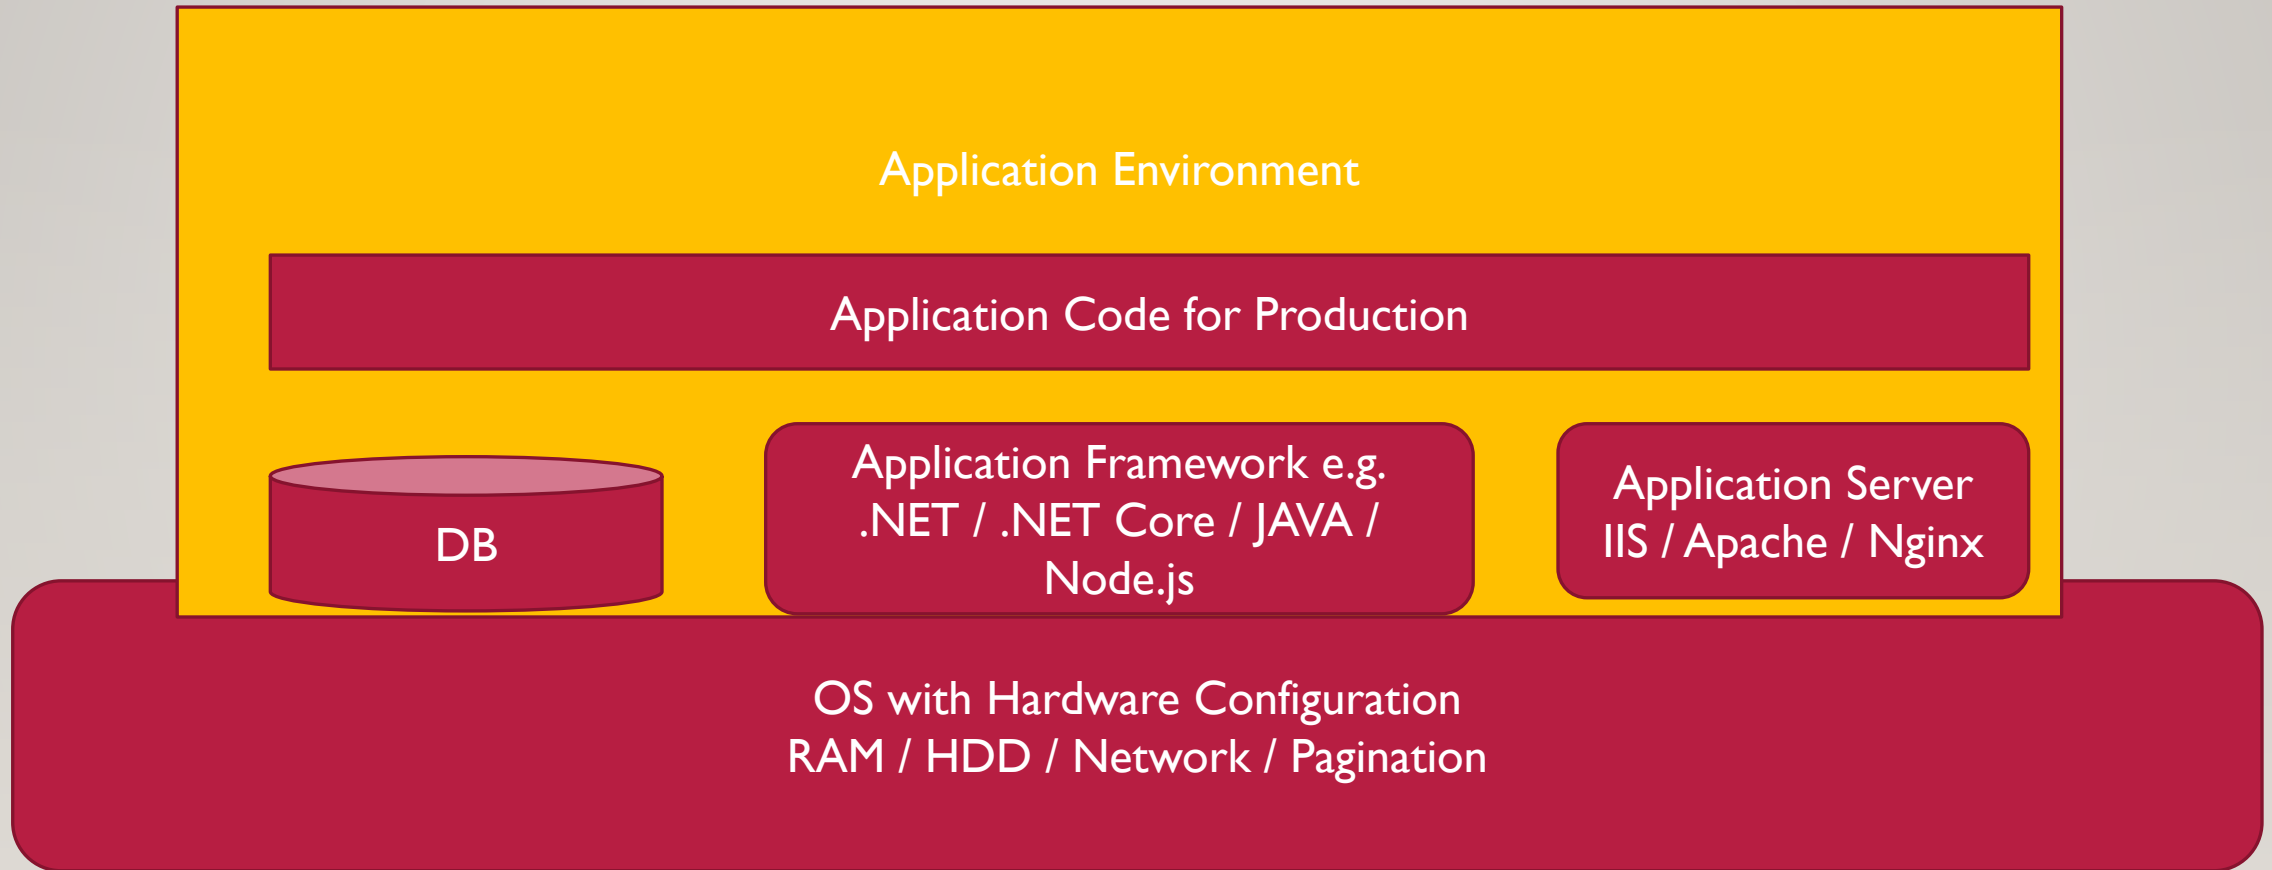
PROGRAMMING WITH MICROSERVICES

- Suitable Technologies for Modern Web Apps
 - ASP.NET Core
 - Node.js
 - Java
- Infrastructure Services
 - Rabbit MQ
 - MassTransit
 - Azure Service Bus
 - Apache Kafka
 - Other messaging Technologies
- Hosting and Management
 - Docker
 - Kubernetes
 - Service Fabric

DOCKER

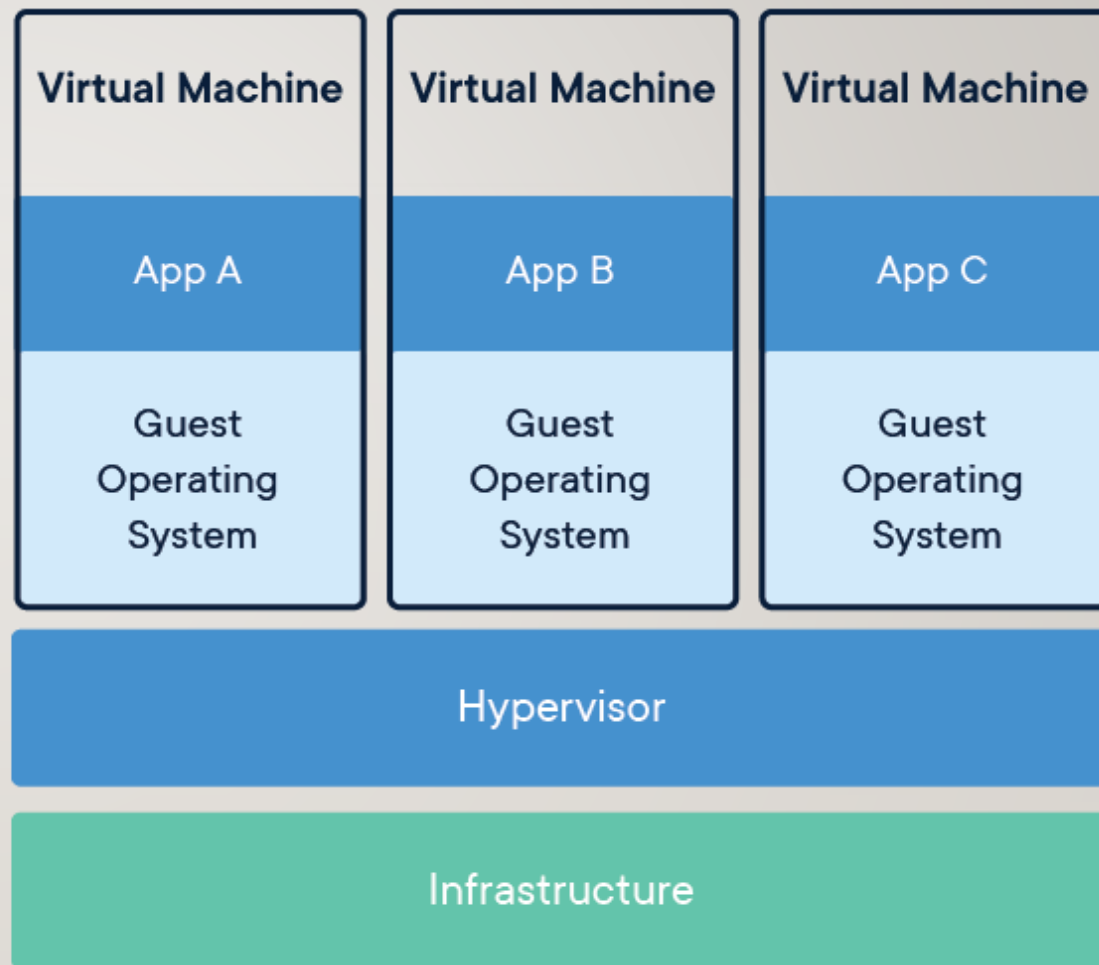
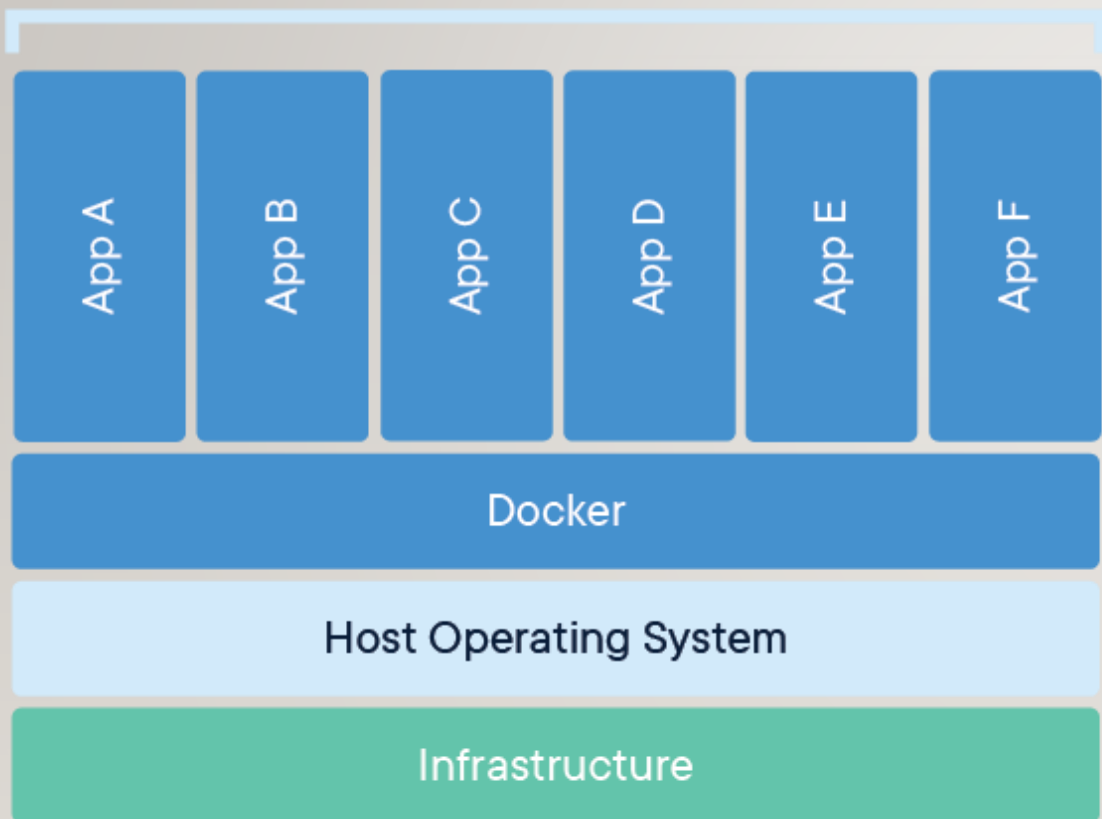


Virtual Env. For Customers



DOCKER

Containerized Applications



```
FROM mcr.microsoft.com/dotnet/core/aspnet:3.1-buster-slim AS base
```

Image from Docker Hub

```
WORKDIR /app
```

Folder on image for app files

```
EXPOSE 80
```

```
EXPOSE 443
```

Port of container to access the app.

```
FROM mcr.microsoft.com/dotnet/core/sdk:3.1-buster AS build
```

```
WORKDIR /src
```

```
COPY Core_NewServie/Core_NewServie.csproj Core_NewServie/
```

```
RUN dotnet restore "Core_NewServie/Core_NewServie.csproj"
```

```
COPY . .
```

```
WORKDIR "/src/Core_NewServie"
```

```
RUN dotnet build "Core_NewServie.csproj" -c Release -o /app/build
```

ASP.NET Core App build is created, create src folder for copying build file from dependencies specified in csproj file. Create a release build

```
FROM build AS publish
```

```
RUN dotnet publish "Core_NewServie.csproj" -c Release -o /app/publish
```

Create a public profile in Docker image, means download all nuget packages (internal and external dependencies) and public application image

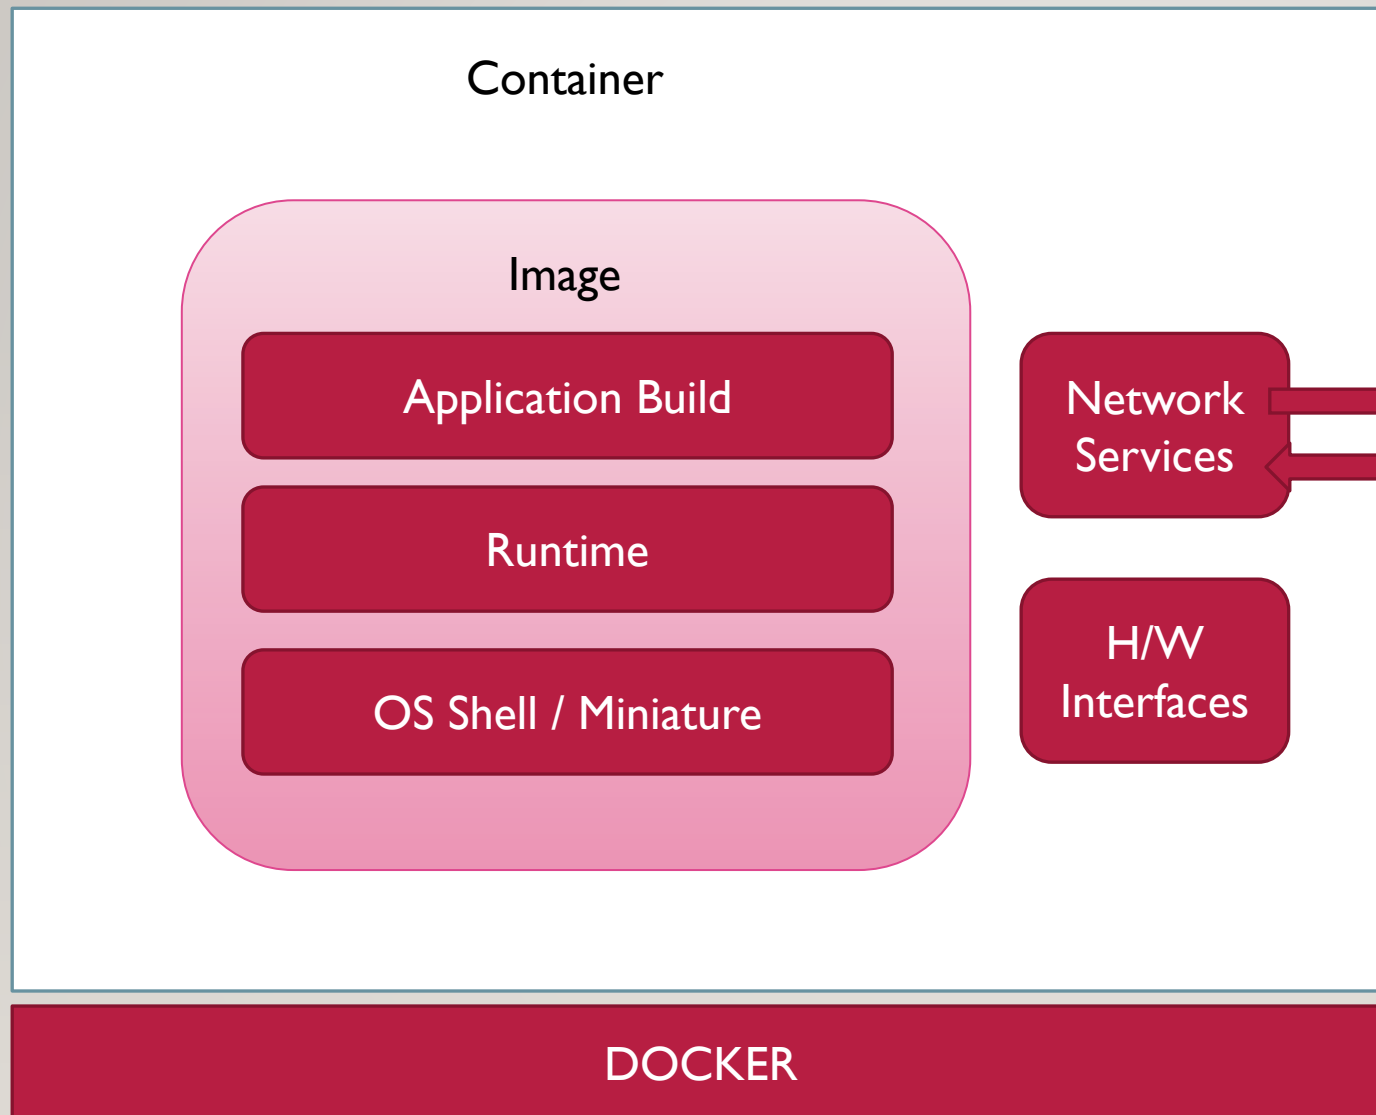
```
FROM base AS final
```

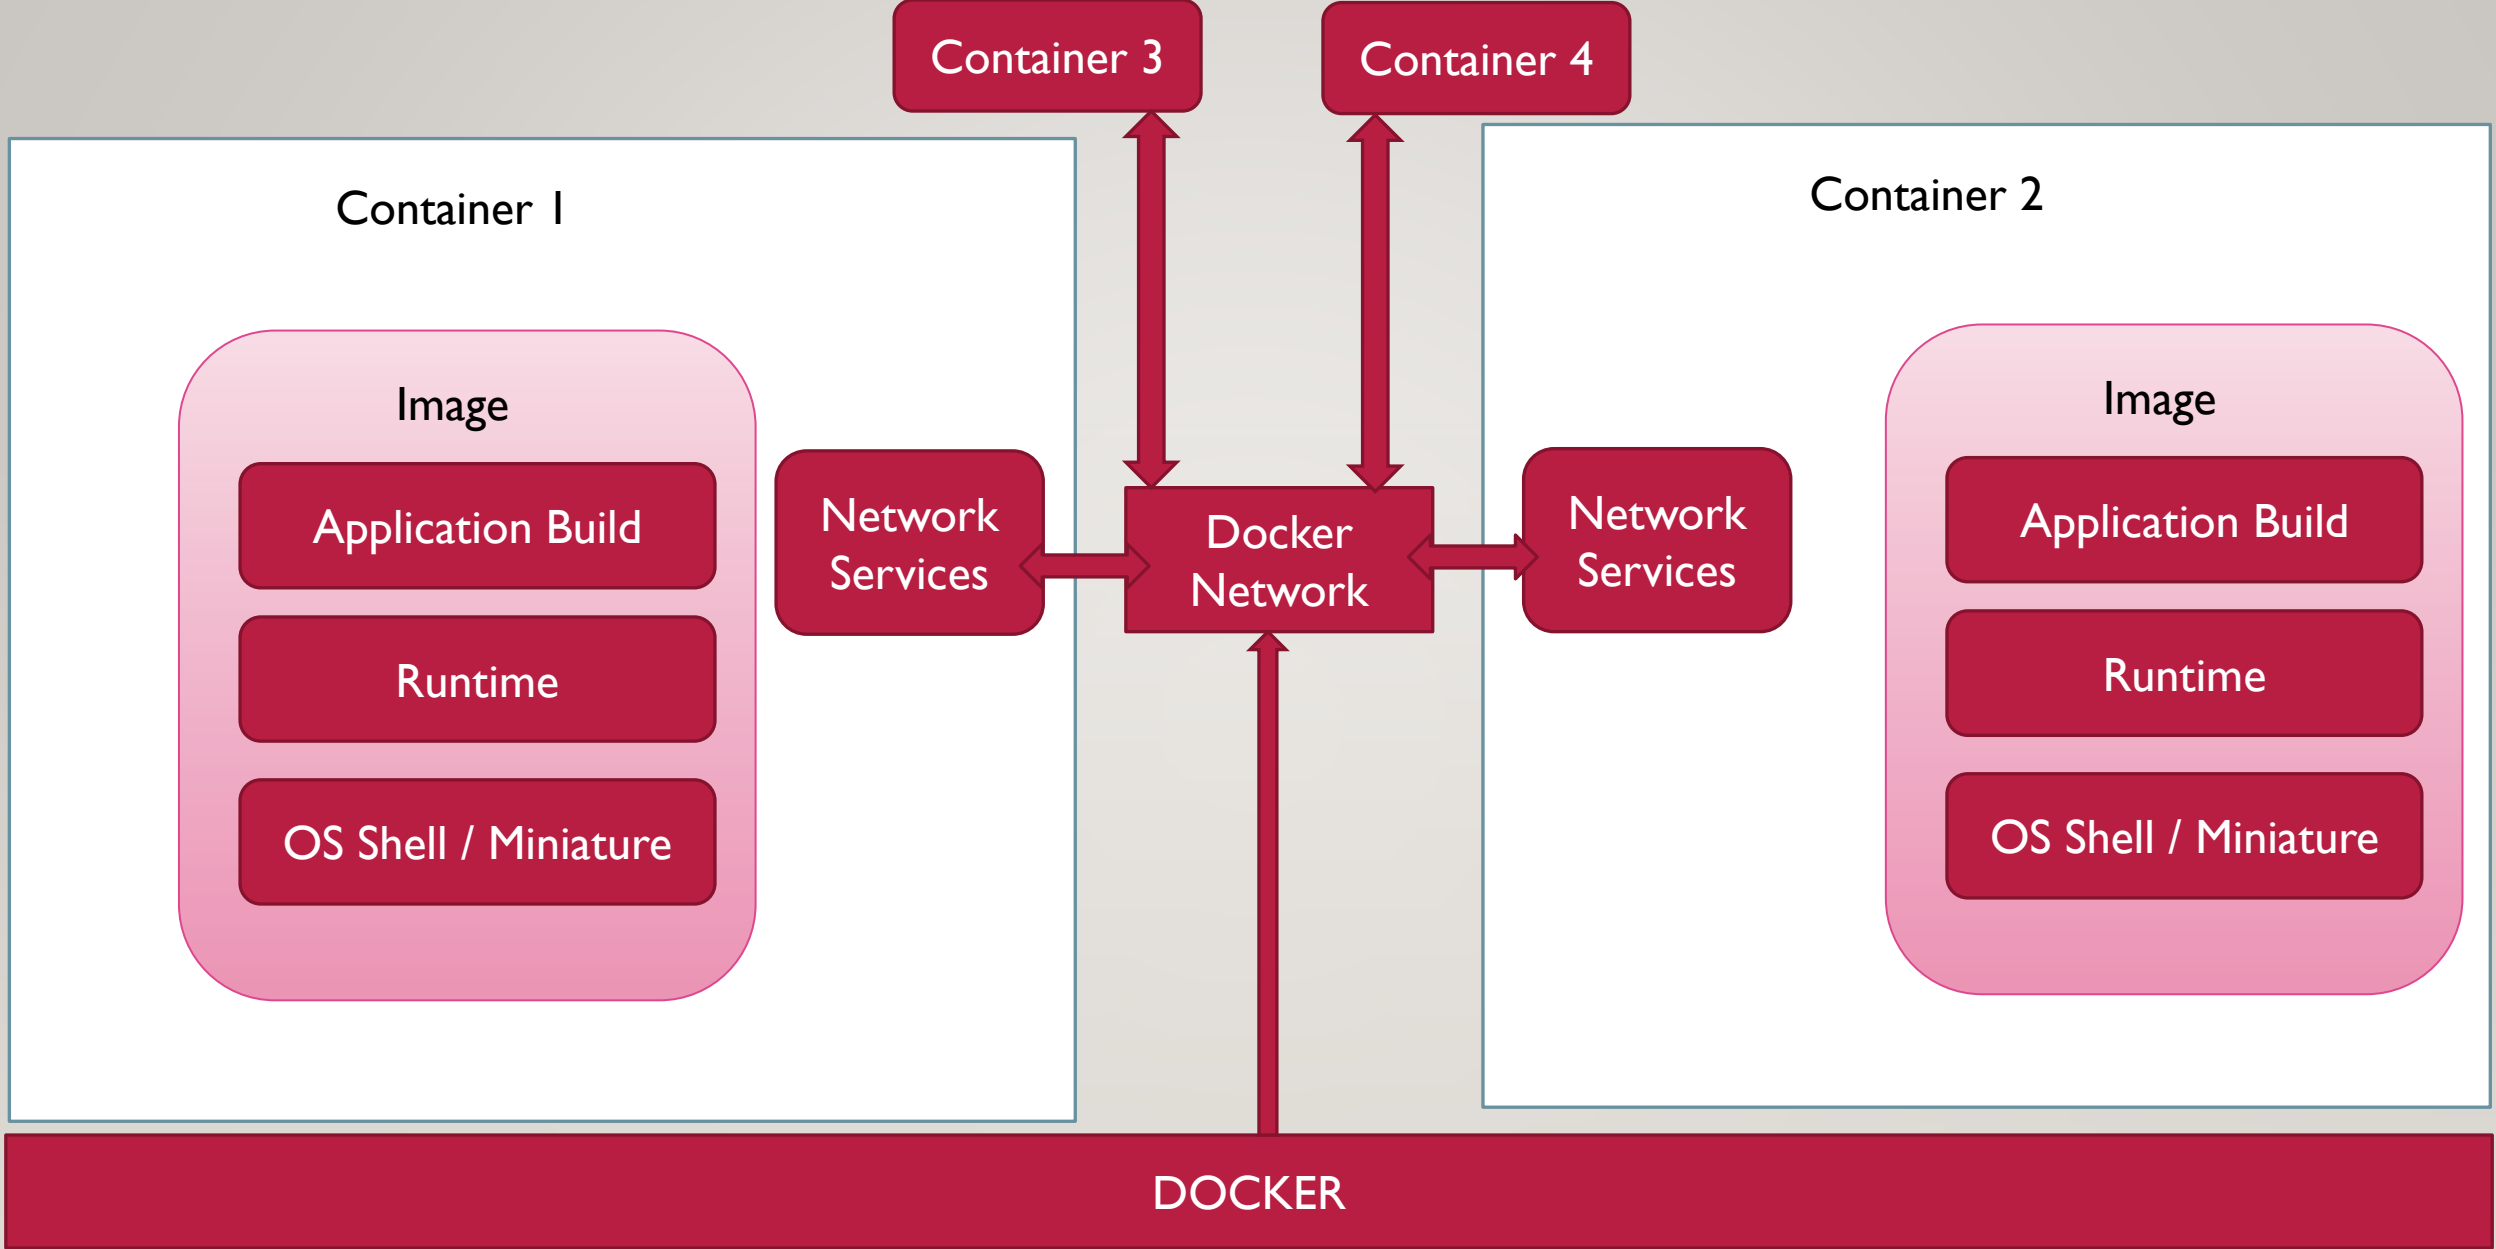
```
WORKDIR /app
```

```
COPY --from=publish /app/publish .
```

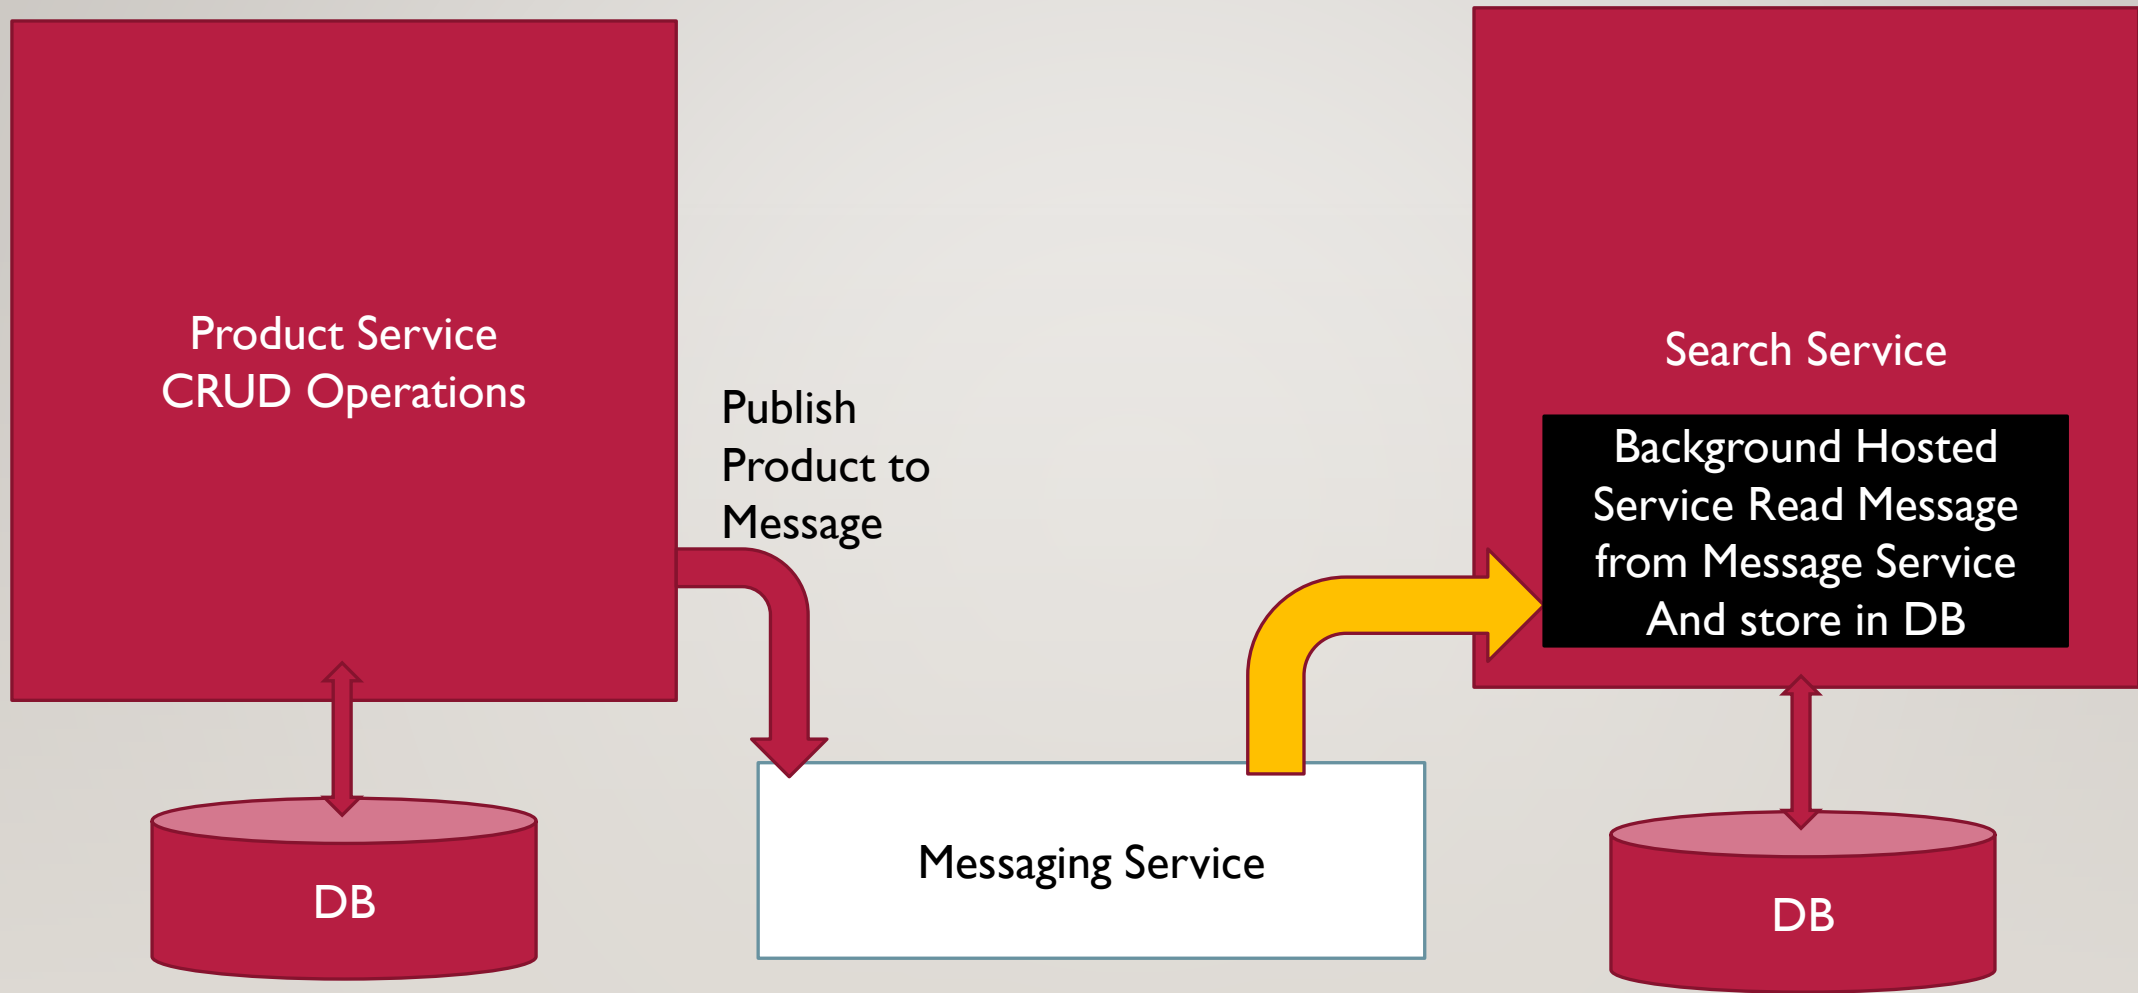
```
ENTRYPOINT ["dotnet", "Core_NewServie.dll"]
```

Image is successfully built and ASP.NET Core Runtime is hosting the application on Port (80/443) and entry-point is set so that users can access this application.





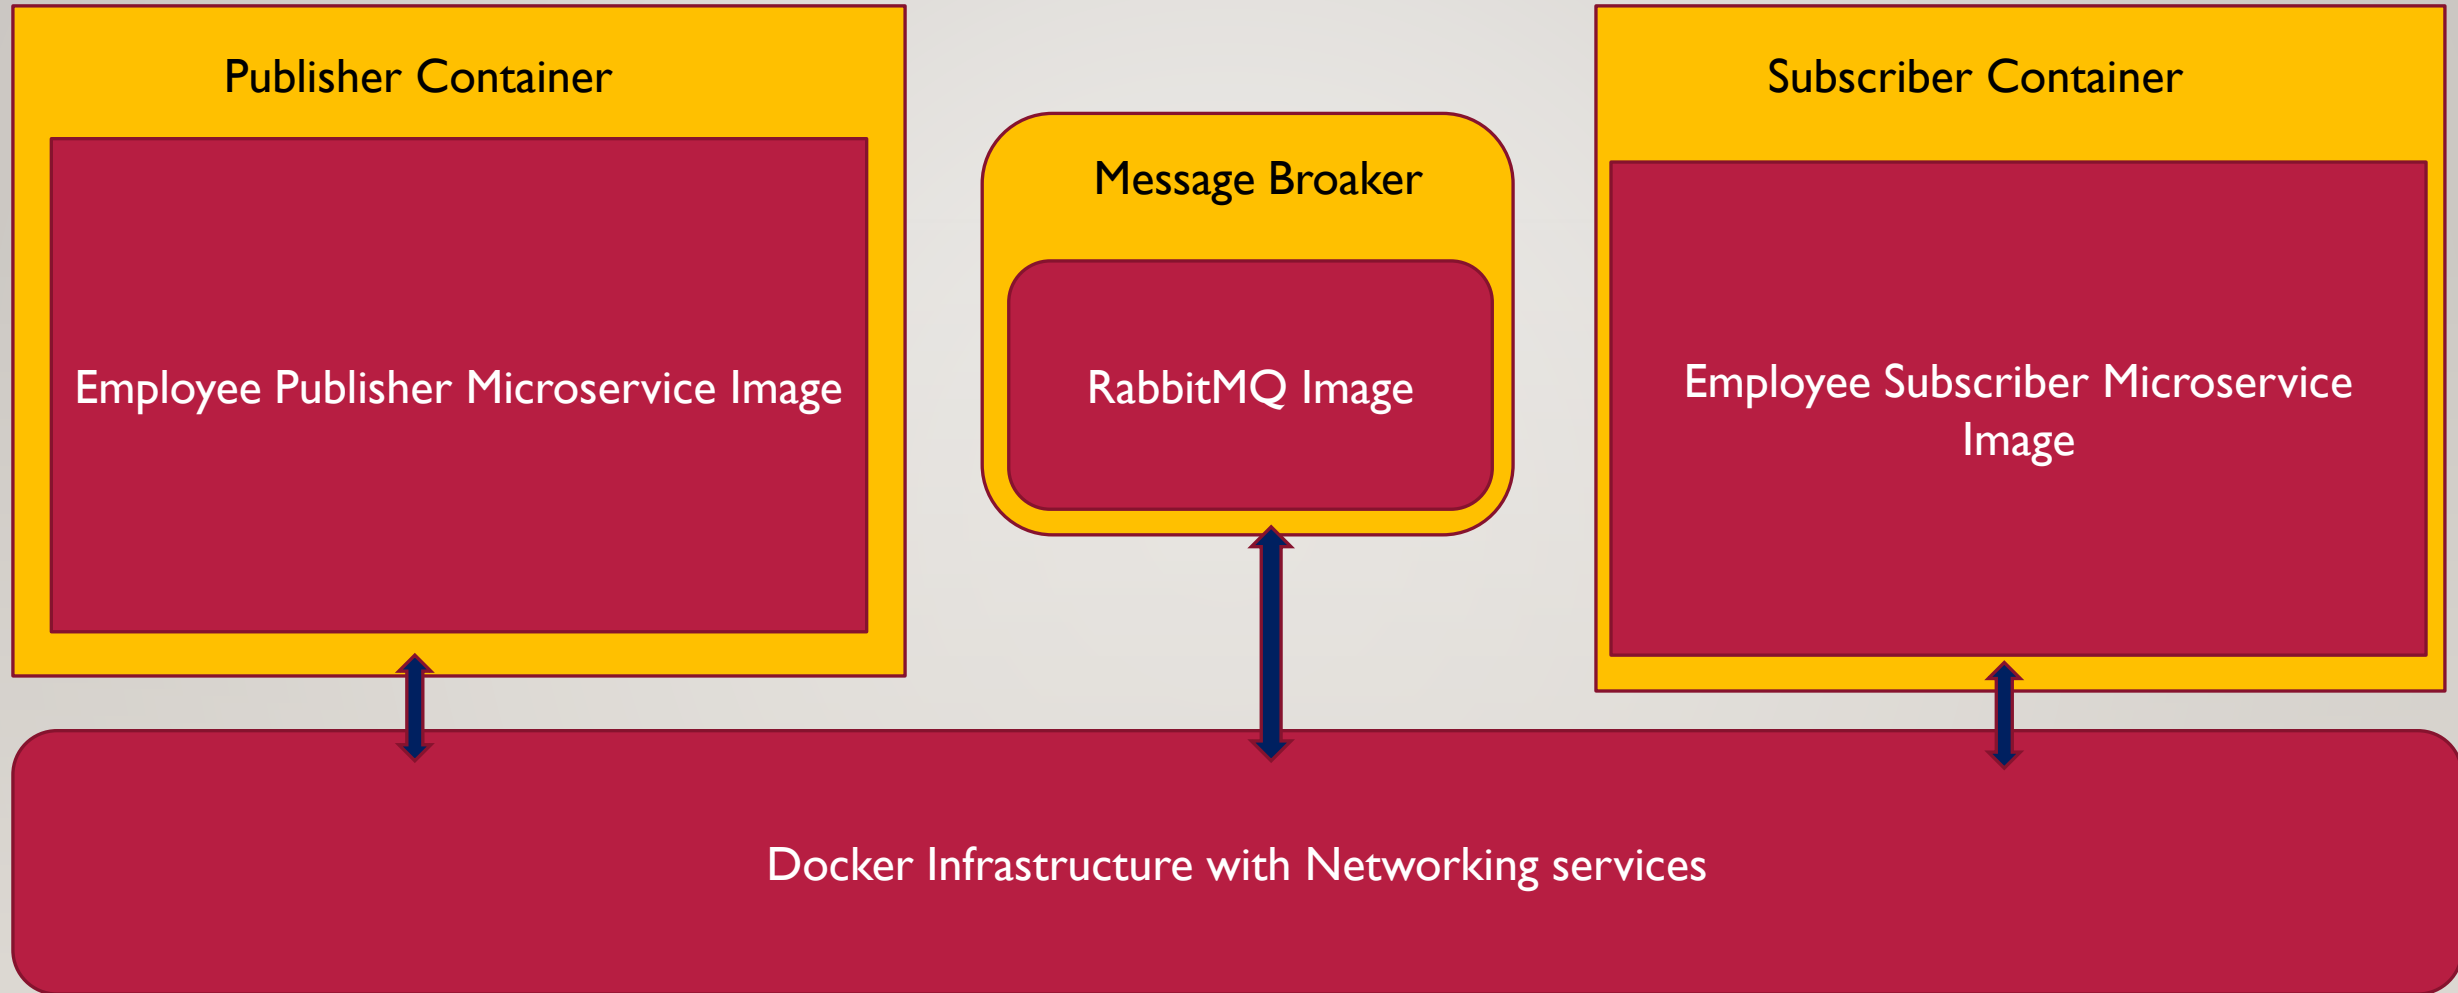
eShopping Microservices App



http://publisher:6001/api/EmployeeAPI

http://U:P@mb:5672

http://subscriber:6002/api/EmployeeAPI



version: "3.6"

services:

rabbitmq:

container_name: emessage-rabbit

ports:

- 5672:5672
- 15672:15672

environment:

- RABBITMQ_DEFAULT_USER=guest
- RABBITMQ_DEFAULT_PASS=guest

image: rabbitmq:3-management

publisherapi:

container_name: publisherapi

ports:

- 9001:80

build: ASP_Net_CorePublisher }

image: netcorepublisher

restart: on-failure

depends_on:

- rabbitmq }

subscriberapi:

container_name: subscriberapi

ports:

- 9002:80

build: ASP_Net_CoreSubscriber

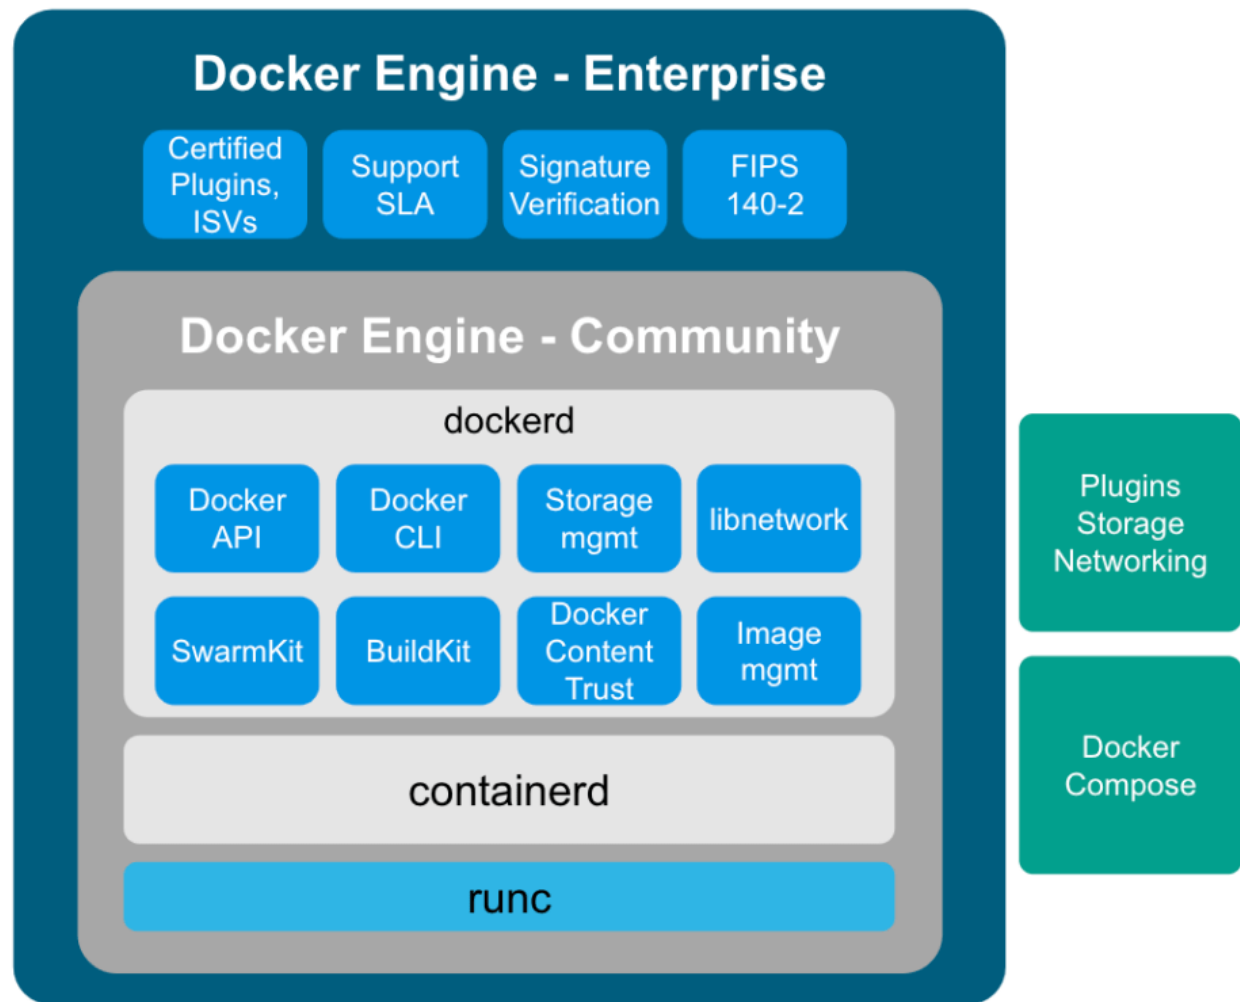
image: netcoresubscriber

restart: on-failure

depends_on:

- rabbitmq

DOCKER



CQRS PATTERN



CQRS

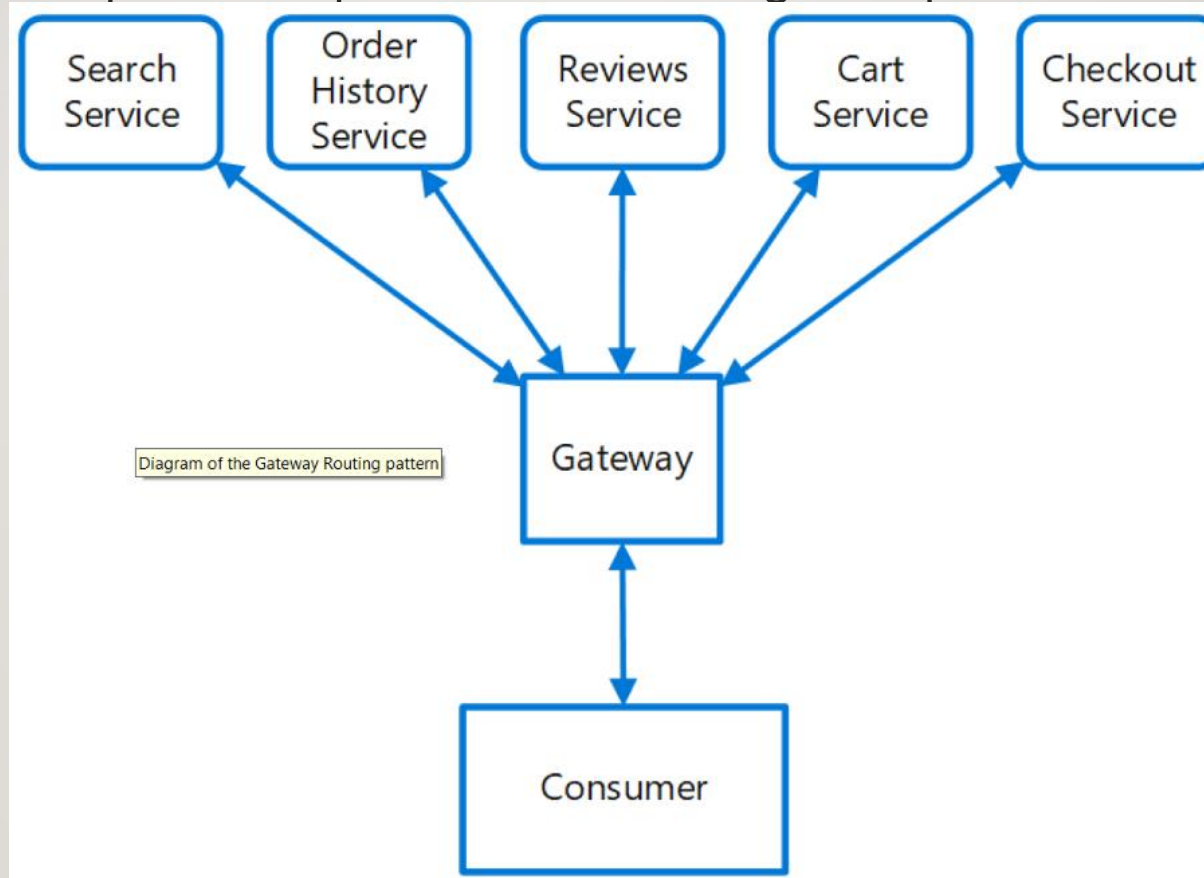
- CQRS stands for Command and Query Responsibility Segregation, a pattern that separates read and update operations for a data store. Implementing CQRS in your application can maximize its performance, scalability, and security. The flexibility created by migrating to CQRS allows a system to better evolve over time and prevents update commands from causing merge conflicts at the domain level.

OCELOT GATEWAY



THE GATEWAY ROUTING PATTERN

- Gateway Routing pattern main objective is Route requests to multiple services using a single endpoint. This pattern is useful when you wish to expose multiple services on a single endpoint and route to the appropriate service based on the request.



Client mobile app



JSON

Back end

Client SPA Web app



JSON

JavaScript/Angular.js

Traditional Web app



HTML

HTML

API Gateway



ASP.NET Core
WebHost

Microservice 1

Web API



Microservice 2

Web API



Microservice 3

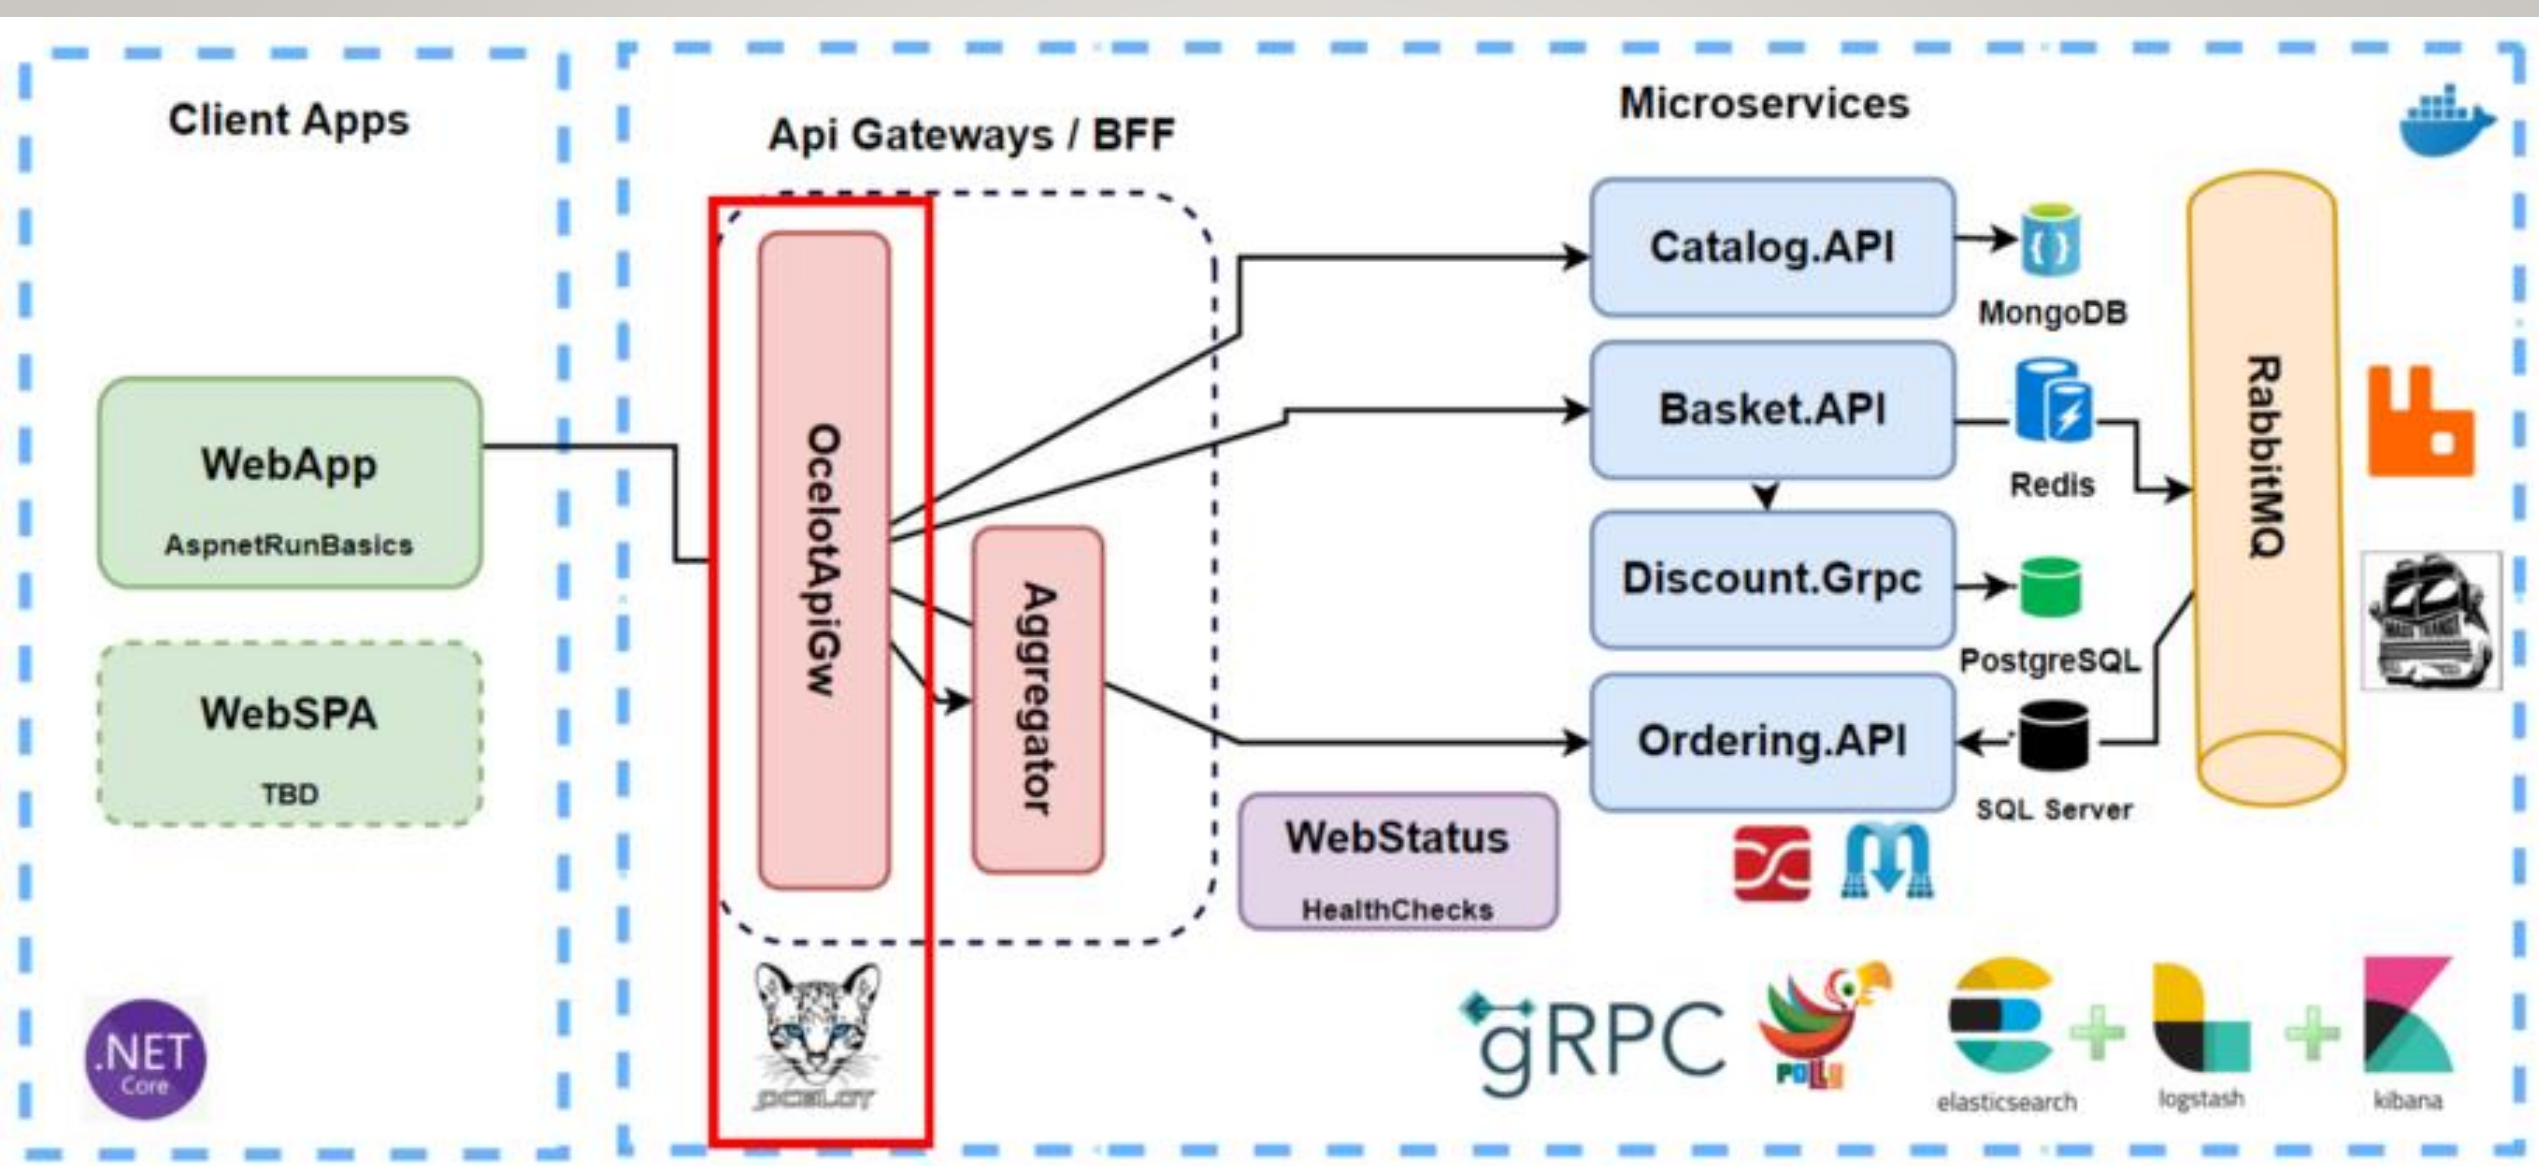
Web API



Client WebApp MVC

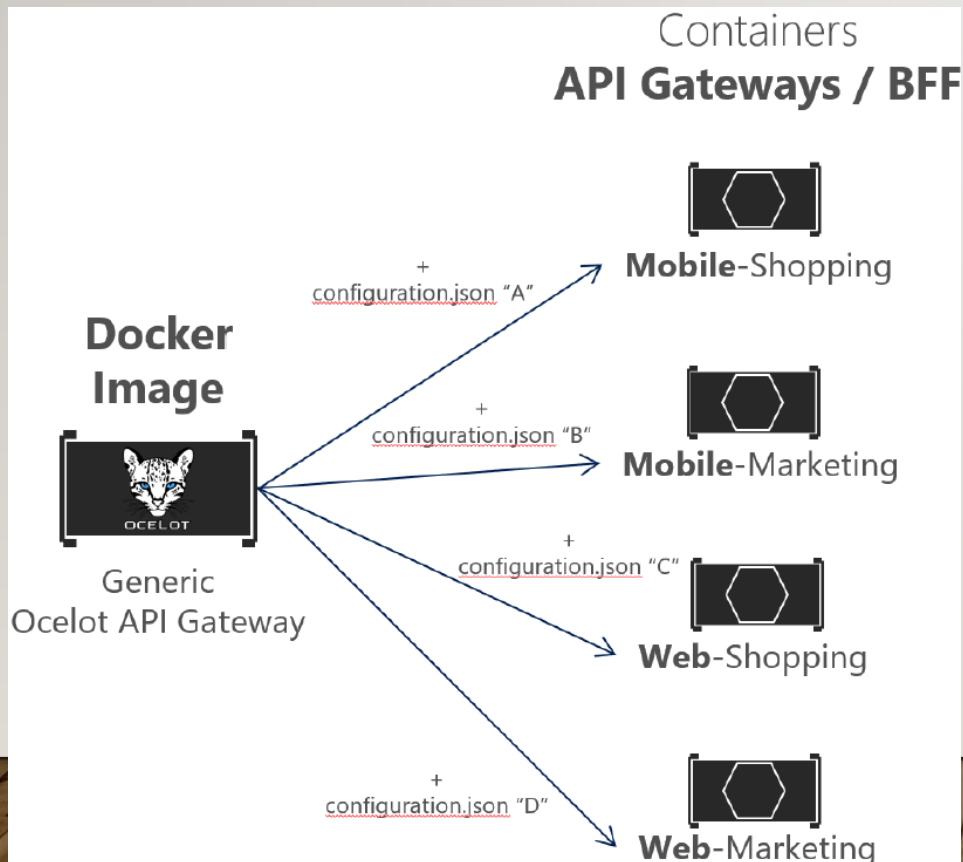
ASP.NET Core MVC
container





OCELOT API GATEWAY


- Ocelot is basically a set of middlewares that you can apply in a specific order.
- Ocelot is a lightweight API Gateway, recommended for simpler approaches. Ocelot is an Open Source .NET Core-based API Gateway especially made for microservices architectures that need unified points of entry into their systems. It's lightweight, fast, and scalable and provides routing and authentication among many other features.



MICROSERVICE RESILIENCY



RESILIENCY



Identifying Failure Points

If you have been troubleshooting your application and individual service performance, you have already likely identified a few services that either receive or send a lot of requests.

Optimizing those requests is important and can help to prolong availability. But, given a high enough load, the services sending or receiving those requests are likely failure points for your application.

For enterprise microservices applications like Uber, where engineers are using thousands upon thousands of microservices, tracing requests across these services can be hopelessly complex – with traces that have hundreds of thousands of spans.

CIRCUIT BREAKER PATTERN

- Microservices applications often rely on remote resources, like third-party services, as a key component of their program.
- But what happens when one of those remote resources times out upon request? Does your microservice continue calling that resource in an endless loop until it fulfills that request?
- What happens when multiple services are requesting that same remote resource?

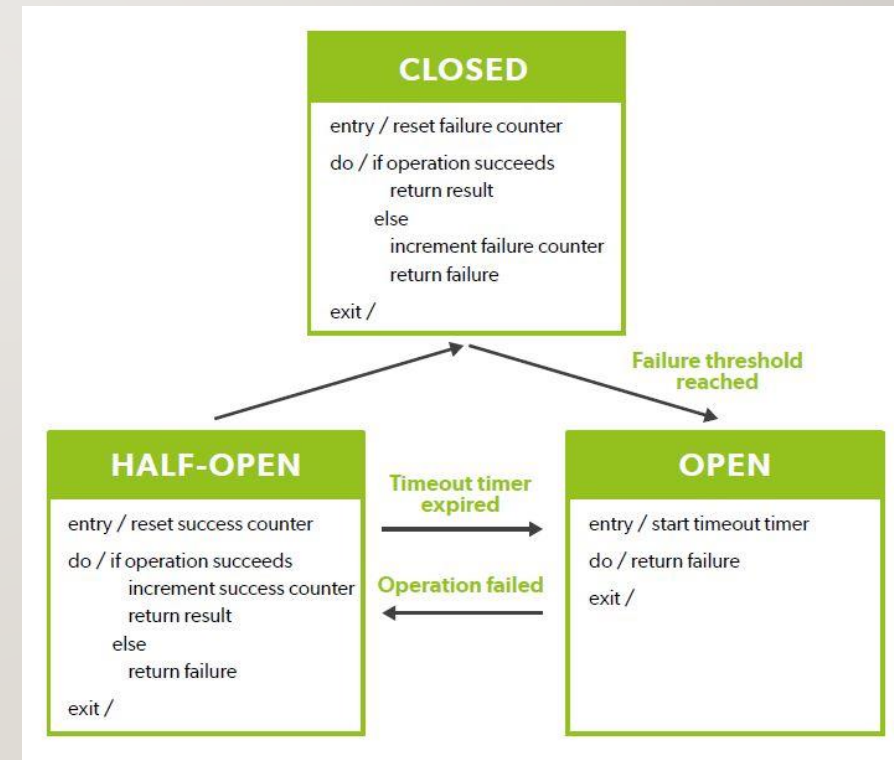
•

What Is the Circuit Breaker Pattern?

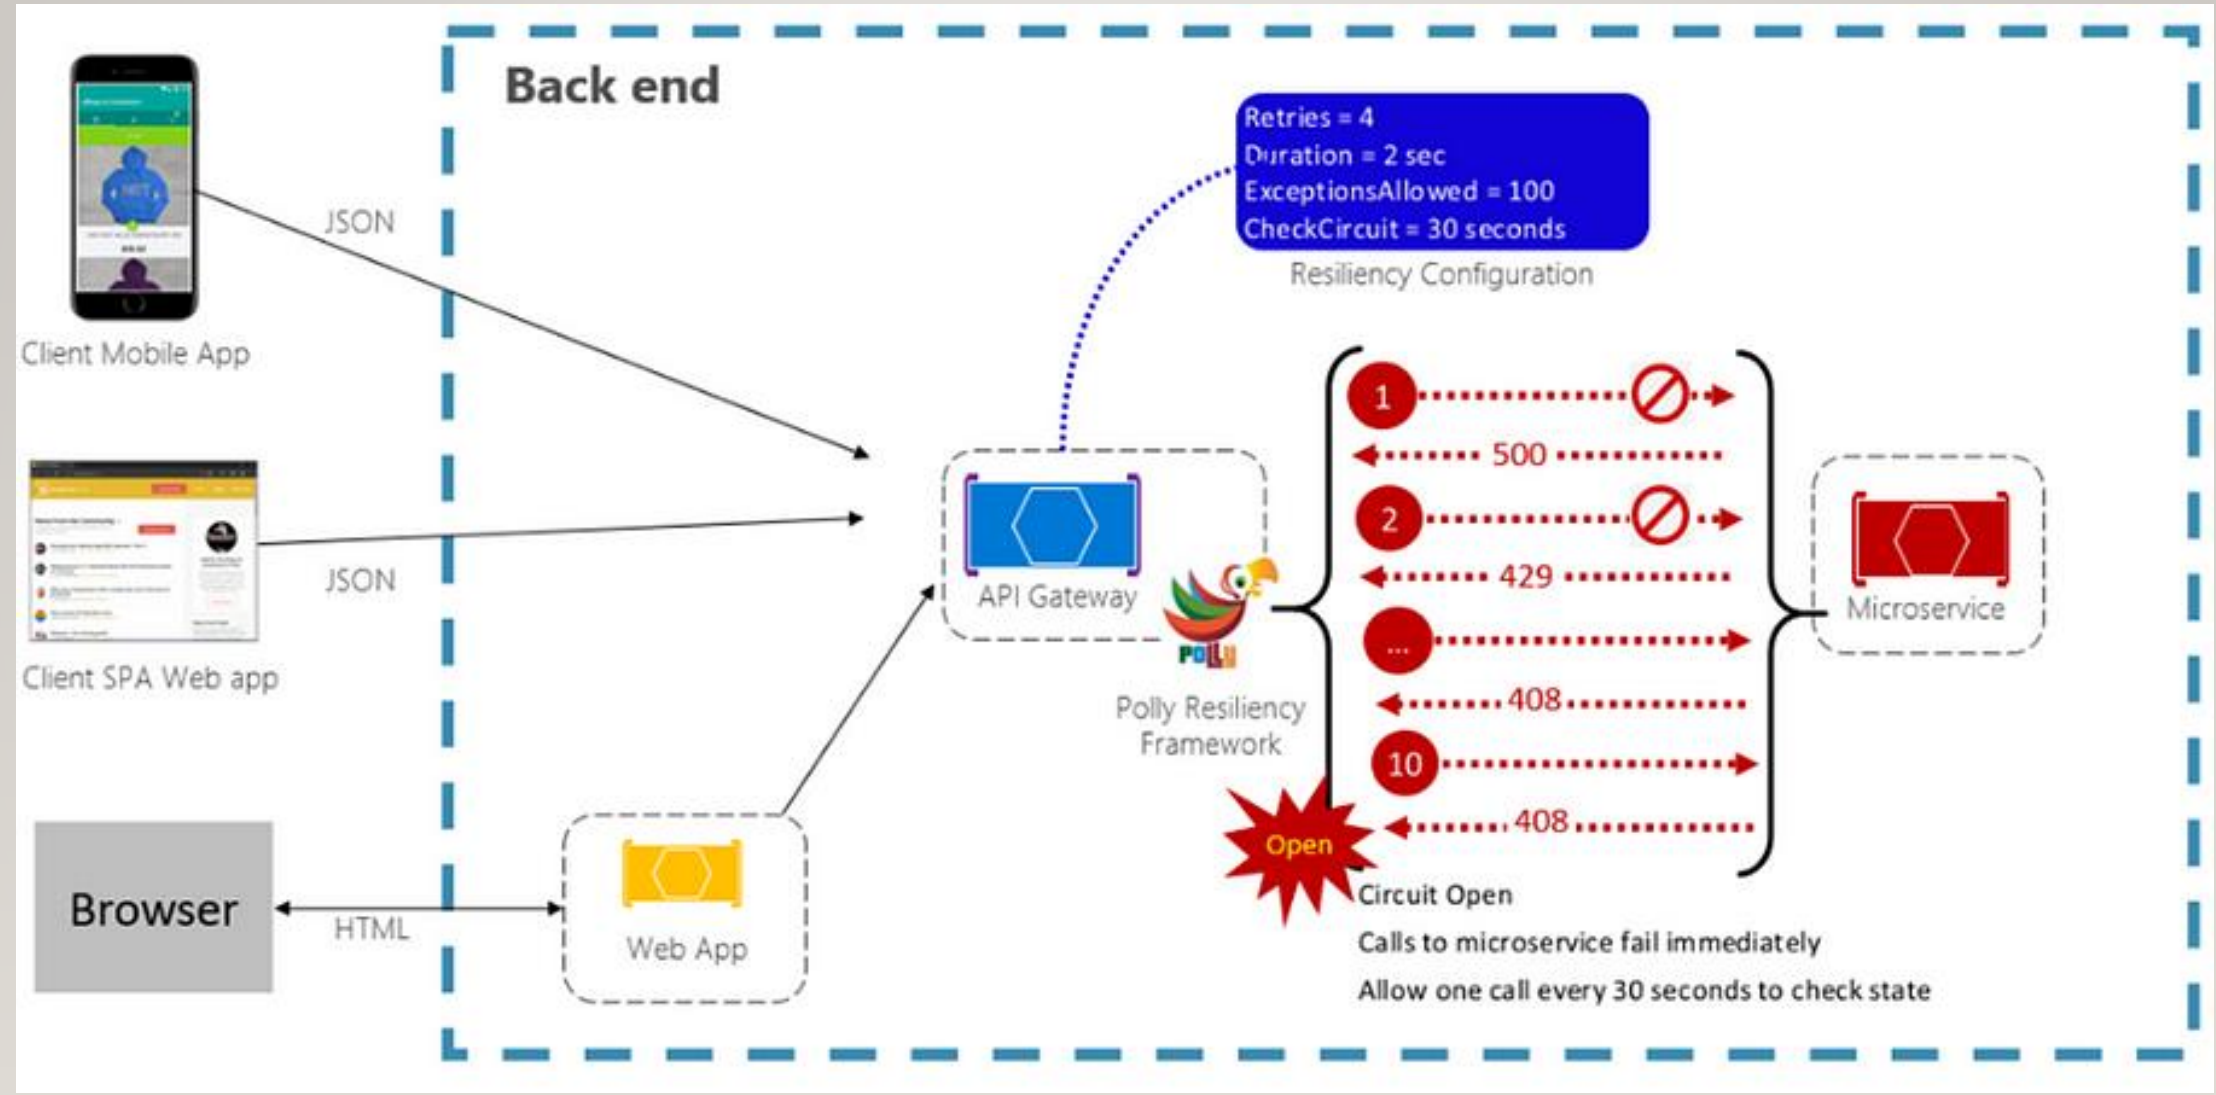
- **The circuit breaker pattern is an application resiliency pattern used to limit the number of requests to a service based on configured thresholds — helping to prevent the service from being overloaded.**

WHEN TO USE THE CIRCUIT BREAKER MICROSERVICES PATTERN

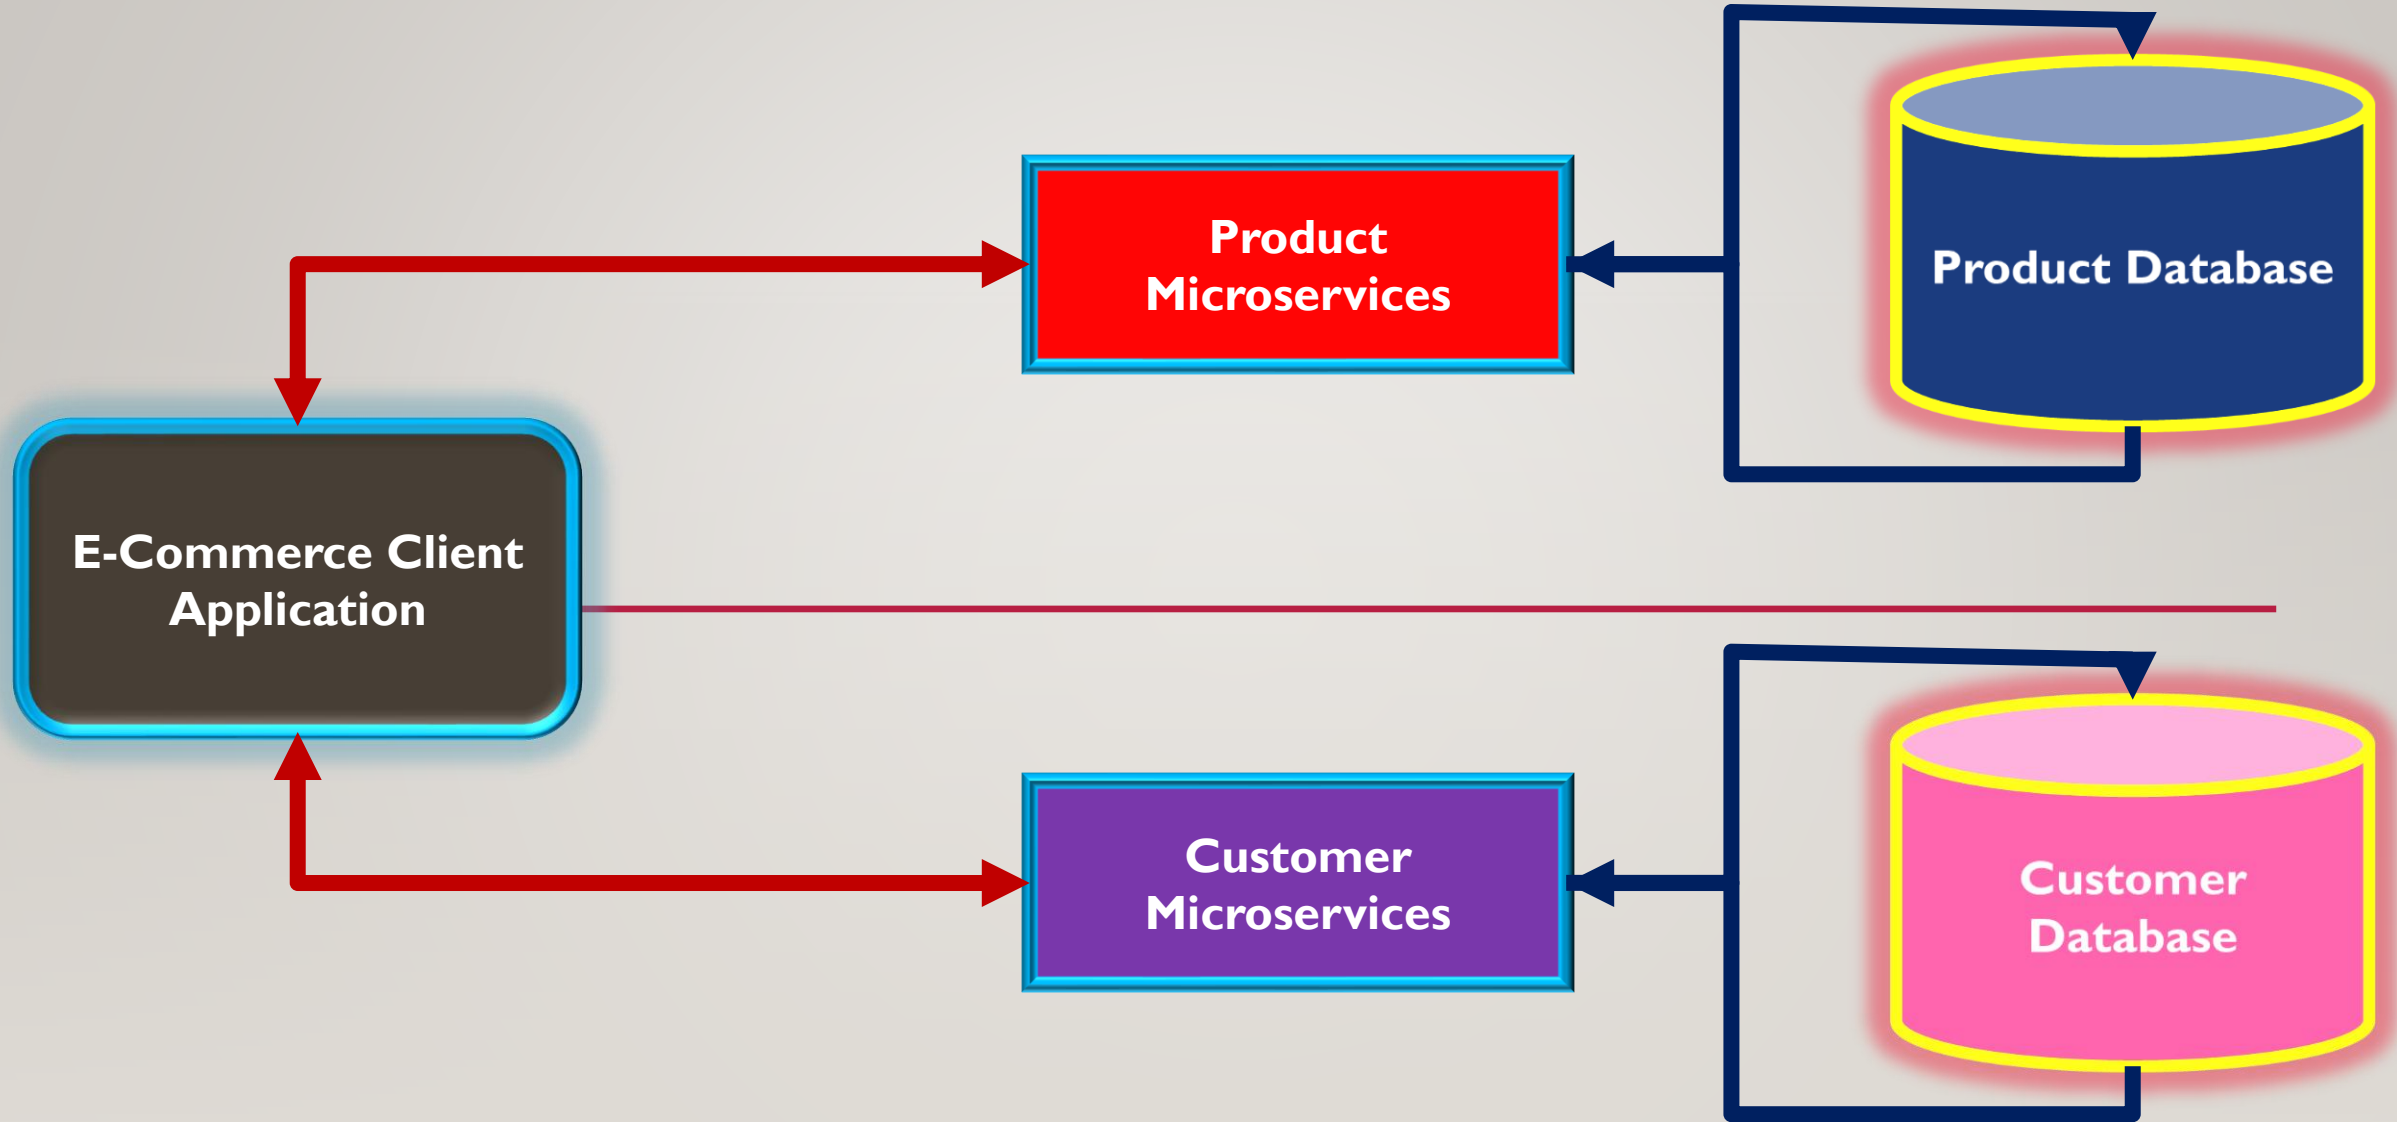
- The circuit breaker pattern can add significant overhead to your application, so it's best to use sparingly, and only in cases where you're accessing a remote service or shared resource prone to failure.
- If you use a service mesh like Istio, it's easy to experiment with various resiliency patterns, including Circuit Breaker, to see which techniques work best for each service.



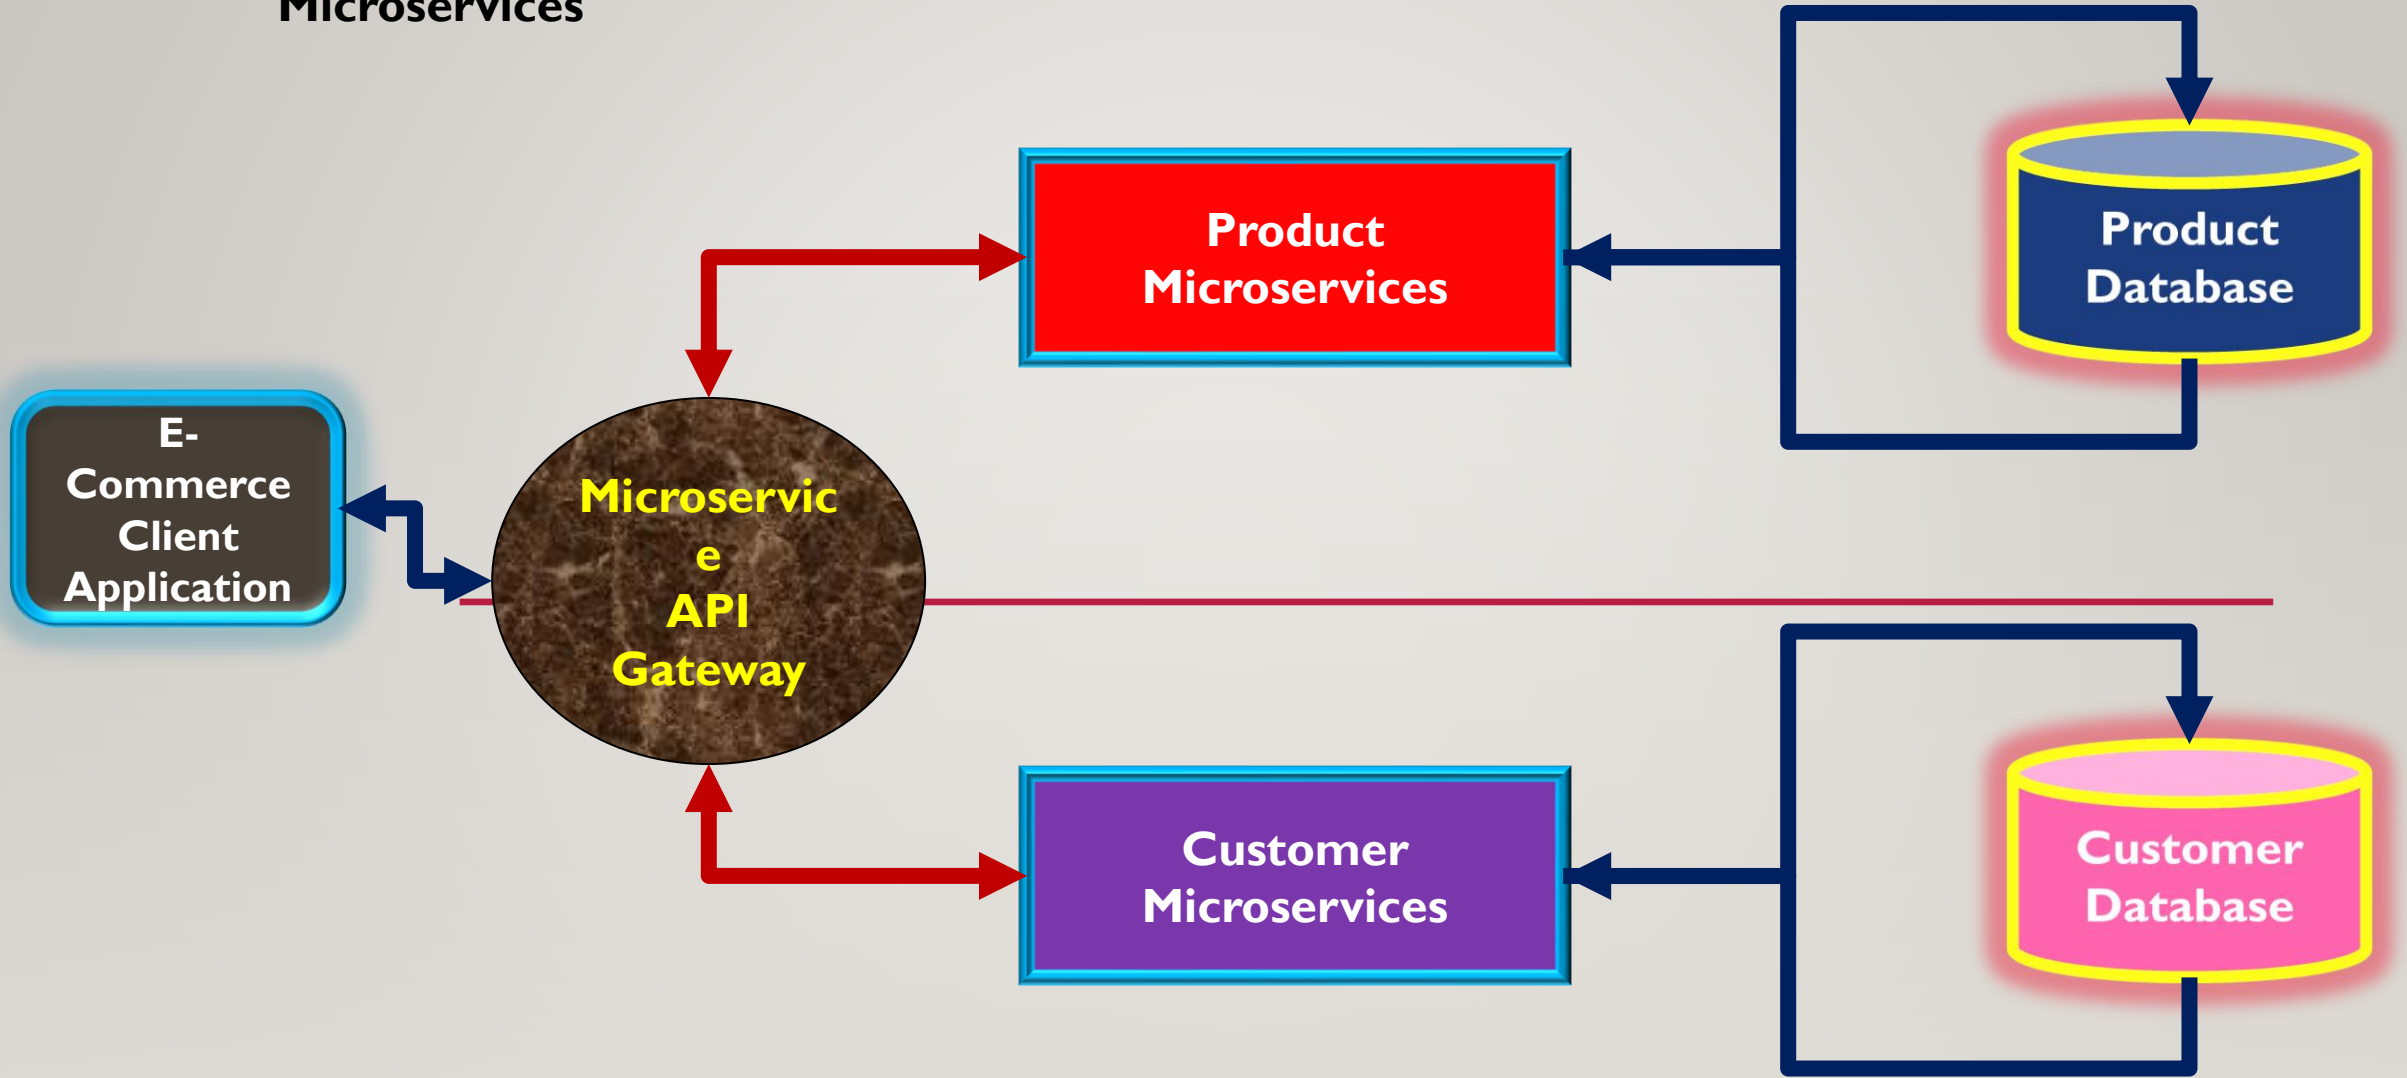
RESILIENC



Direct Access of the Microservices



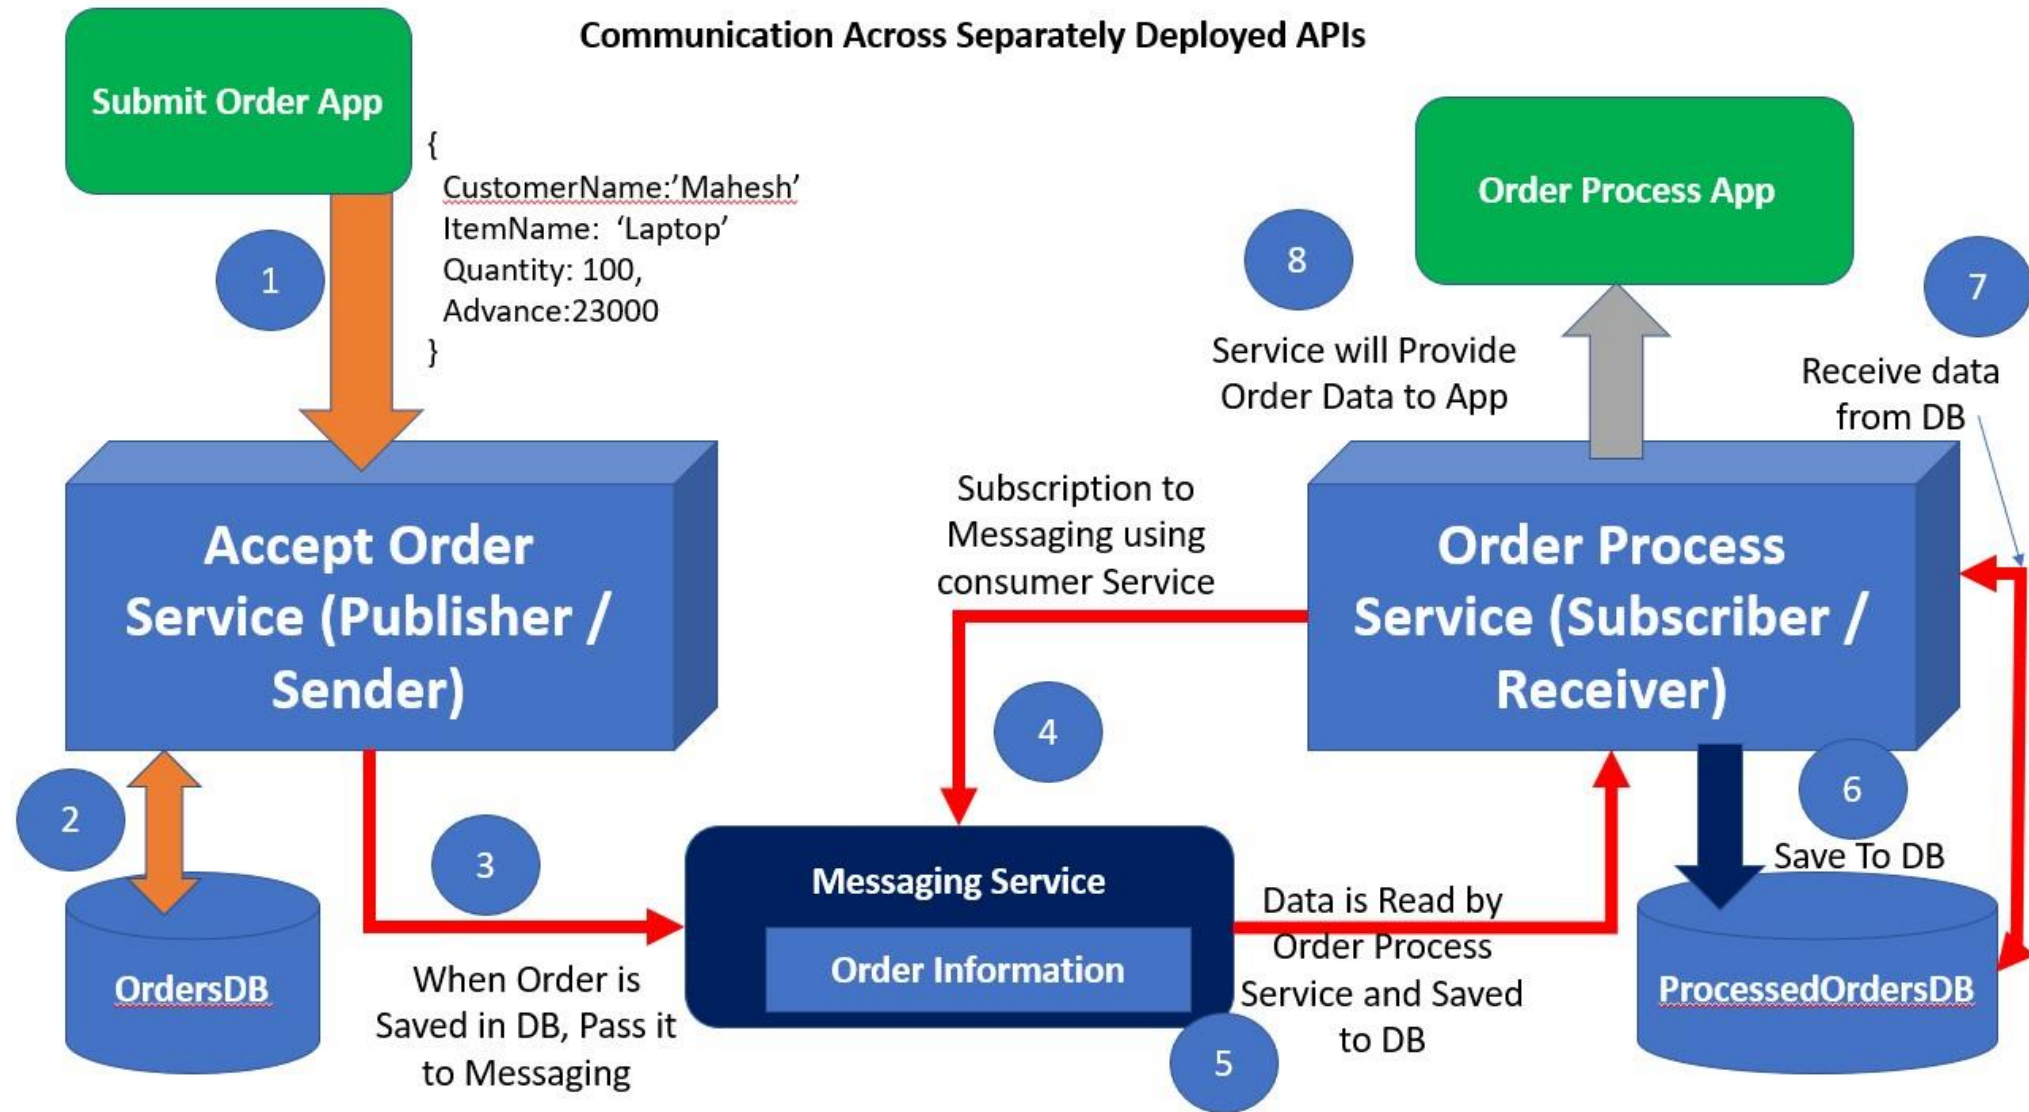
Resilience Design of the Microservices



MICROSERVICES COMMUNICATION



Communication Across Separately Deployed APIs

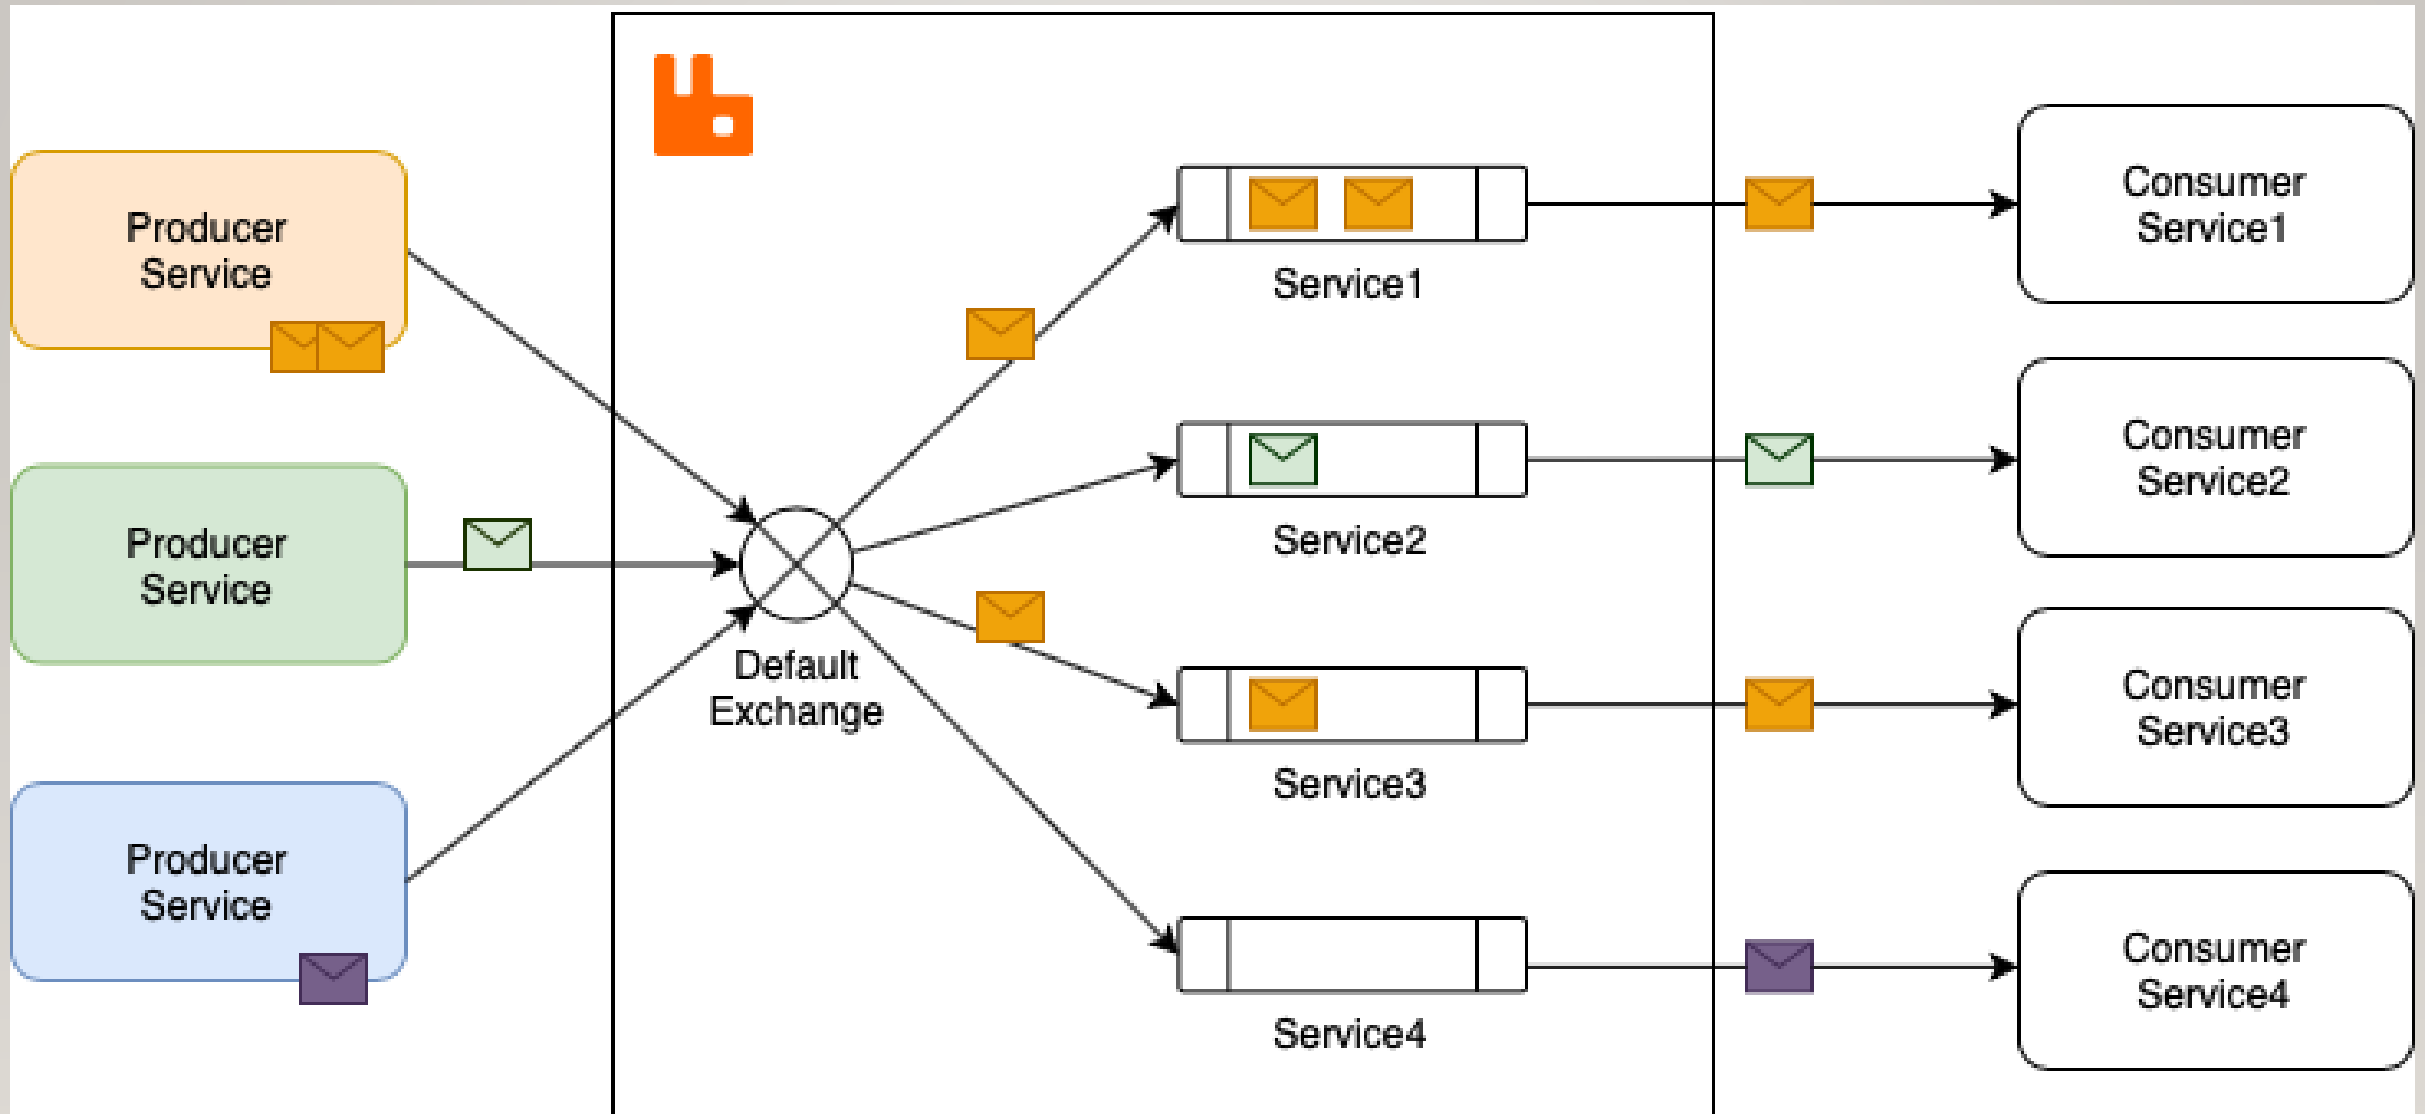


RABBITMQ

- RabbitMQ is the most widely deployed open source message broker.
- RabbitMQ is lightweight and easy to deploy on premises and in the cloud. It supports multiple messaging protocols. RabbitMQ can be deployed in distributed and federated configurations to meet high-scale, high-availability requirements.
- It handles the accepting, storing, and sending of messages between our applications. Using a message broker allows us to build decoupling, performant applications, relying on asynchronous communication between our applications.
- **The durable queue will keep message persisted even after broker restarts**

INSTALLING RABBITMQ

- `docker run -it --rm --name rabbitmq -p 5672:5672 -p 15672:15672 rabbitmq:3.11-management`




WHAT ARE RABBITMQ EXCHANGES?

- When working with RabbitMQ, producers can send messages to a couple of different endpoints:
 - Queues
 - Exchanges
- When a producer sends directly to a queue, this message will be received by all consumers of that queue. But what if we want to selectively send messages to different queues based on metadata found in the message? This is where exchanges come into play.
- An exchange receives messages from producers, and depending on its configuration, will send the message to one or many queues. We must create a **binding**, which will ensure our messages get sent from our exchange to one or many queues.
- We can define exchanges from one of the following types:
 - Direct
 - Topic
 - Headers
 - Fanout

RabbitMQ

Producers



Subscribers



MASS TRANSIT

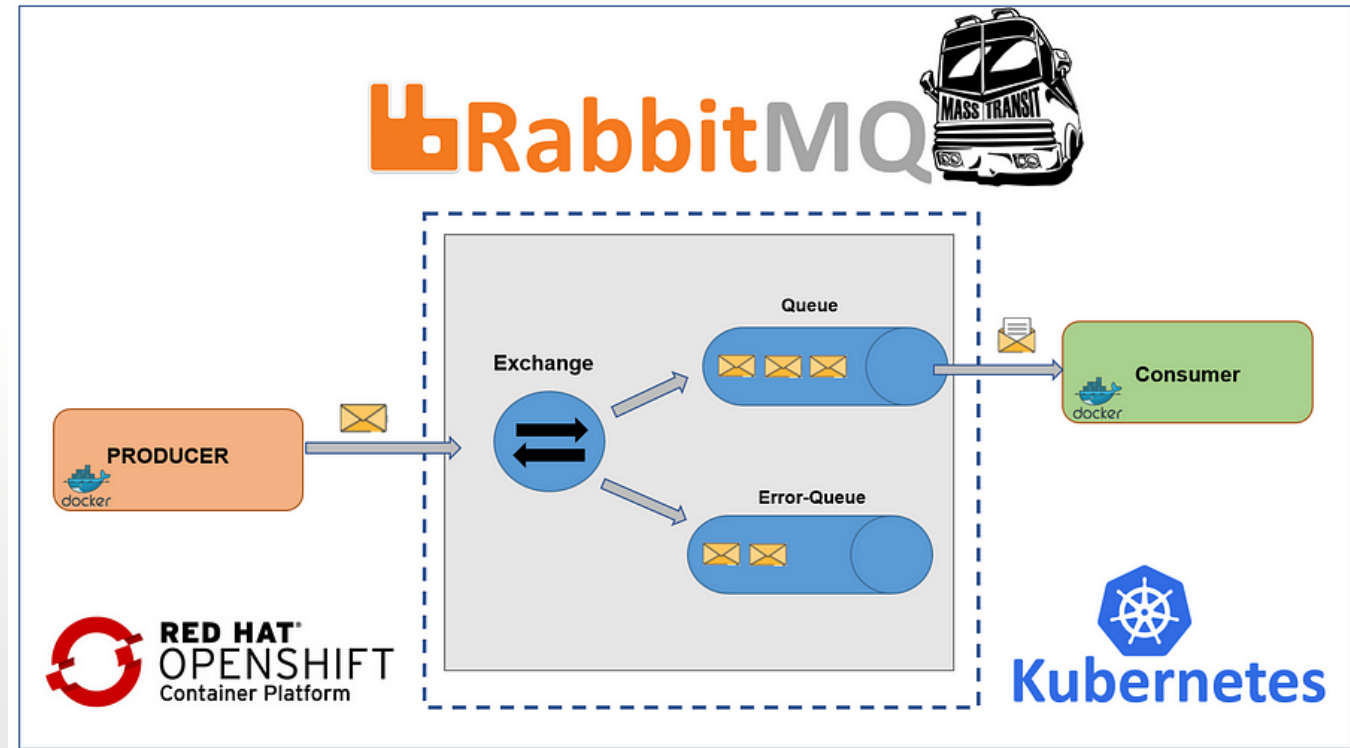


MASSTRANSIT

- **MassTransit** is a free, open-source, distributed application framework for .NET applications.
- **It abstracts away the underlying logic required to work with message brokers, such as RabbitMQ, making it easier to create message-based, loosely coupled applications.**
- **Service Bus**, usually shortened to **Bus**, is the term given to the type of application that handles the movement of messages.
- **Transports** are the different types of message brokers MassTransit works with, including RabbitMQ, InMemory, Azure Service Bus, and more.
- **Message** is a contract, defined *code first* by creating a .NET class or interface.
- **Command** is a type of message, specifically used to tell a service to do something. These message types are **sent** to an endpoint (queue) and will be expressed using a verb-noun sequence.
- **Events** are another message type, signifying that something has happened. Events are **published** to one or multiple consumers and will be expressed using noun-verb (past tense) sequence.

MASSTRANSIT

- **Why Use MassTransit?**
- There are a few benefits to choosing to use a library such as MassTransit, instead of working with the native message broker library.
- Firstly, **by abstracting the underlying message broker logic, we can work with multiple message brokers**, without having to completely rewrite our code.
 - This allows us to work with something such as the InMemory transport when working locally, then when deploying our code, use another transport such as Azure Service Bus or Amazon Simple Queue Service.
- Additionally, when we work with a message-based architecture, **there are a lot of specific patterns we need to be aware of and implement, such as *retry*, *circuit breaker*, *outbox* to name a few.**
- **MassTransit handles all of this for application**, along with many other features such as *exception handling*, *distributed transactions*, and *monitoring*.



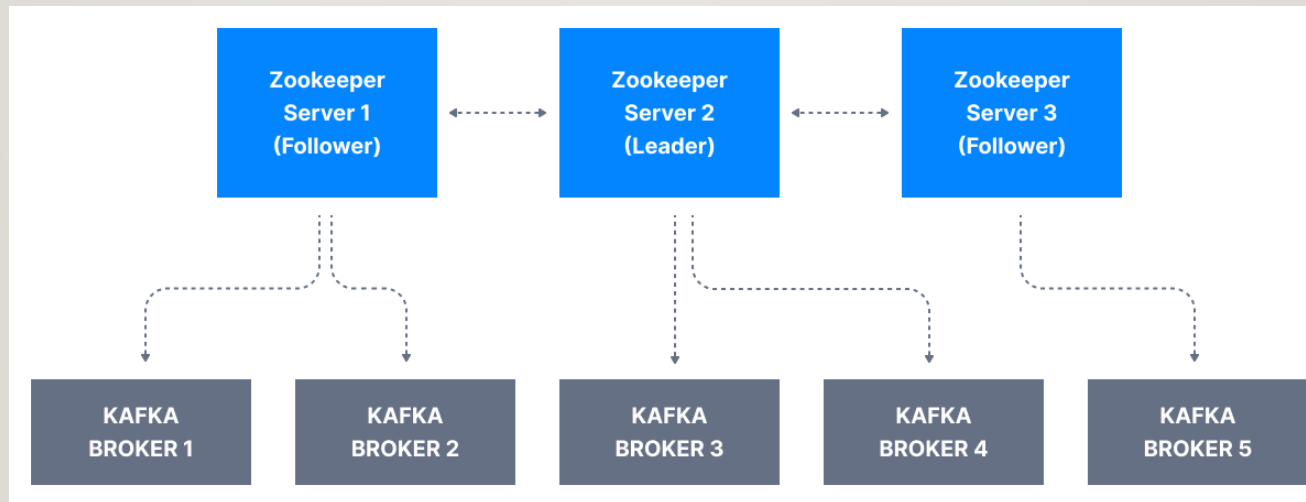
KAFKA

KAFKA

- Apache Kafka is an open-source distributed streaming system used for stream processing
- Apache Kafka is a distributed event store and stream-processing platform. It is an open-source system developed by the Apache Software Foundation written in Java and Scala.
- The project aims to provide a unified, high-throughput, low-latency platform for handling real-time data feeds.


ZOOKEEPER

- Zookeeper is used for metadata management in the Kafka world. For example:
- Zookeeper keeps track of which brokers are part of the Kafka cluster
- Zookeeper is used by Kafka brokers to determine which broker is the leader of a given partition and topic and perform leader elections
- Zookeeper stores configurations for topics and permissions
- Zookeeper sends notifications to Kafka in case of changes (e.g. new topic, broker dies, broker comes up, delete topics, etc....)



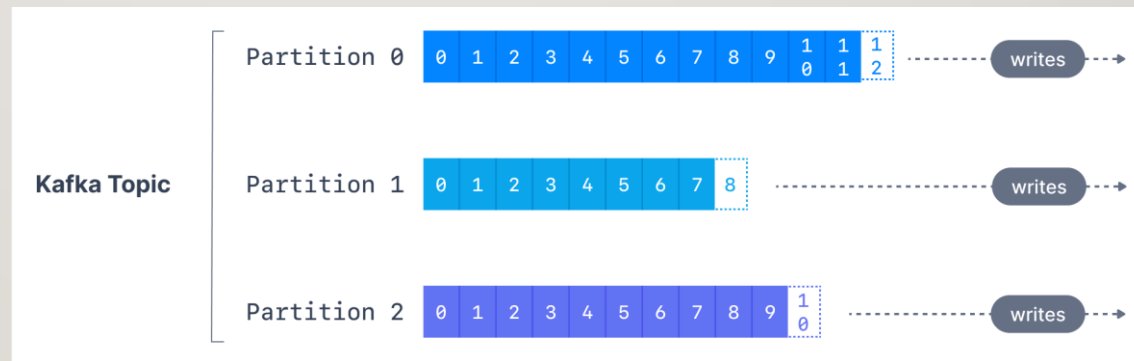
KAFKA TOPICS

- **What is a Kafka Topic?**
- Similar to how databases have tables to organize and segment datasets, Kafka uses the concept of topics to organize related messages.
- A topic is identified by its name. For example, we may have a topic called **logs** that may contain log messages from our application, and another topic called **purchases** that may contain purchase data from our application as it happens.
- Kafka topics can contain any kind of message in any format, and the sequence of all these messages is called a data stream.
- Data in Kafka topics is deleted after one week by default (also called the default message retention period), and this value is configurable. This mechanism of deleting old data ensures a Kafka cluster does not run out of disk space by recycling topics over time.



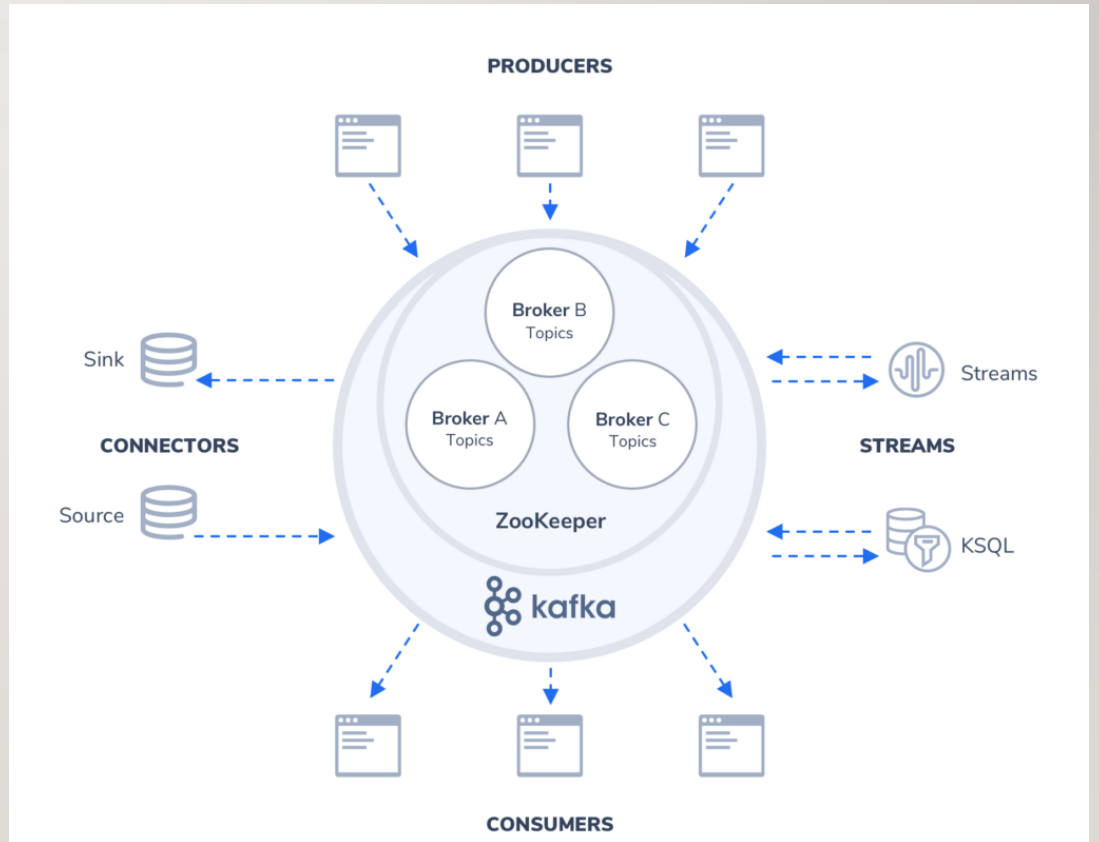
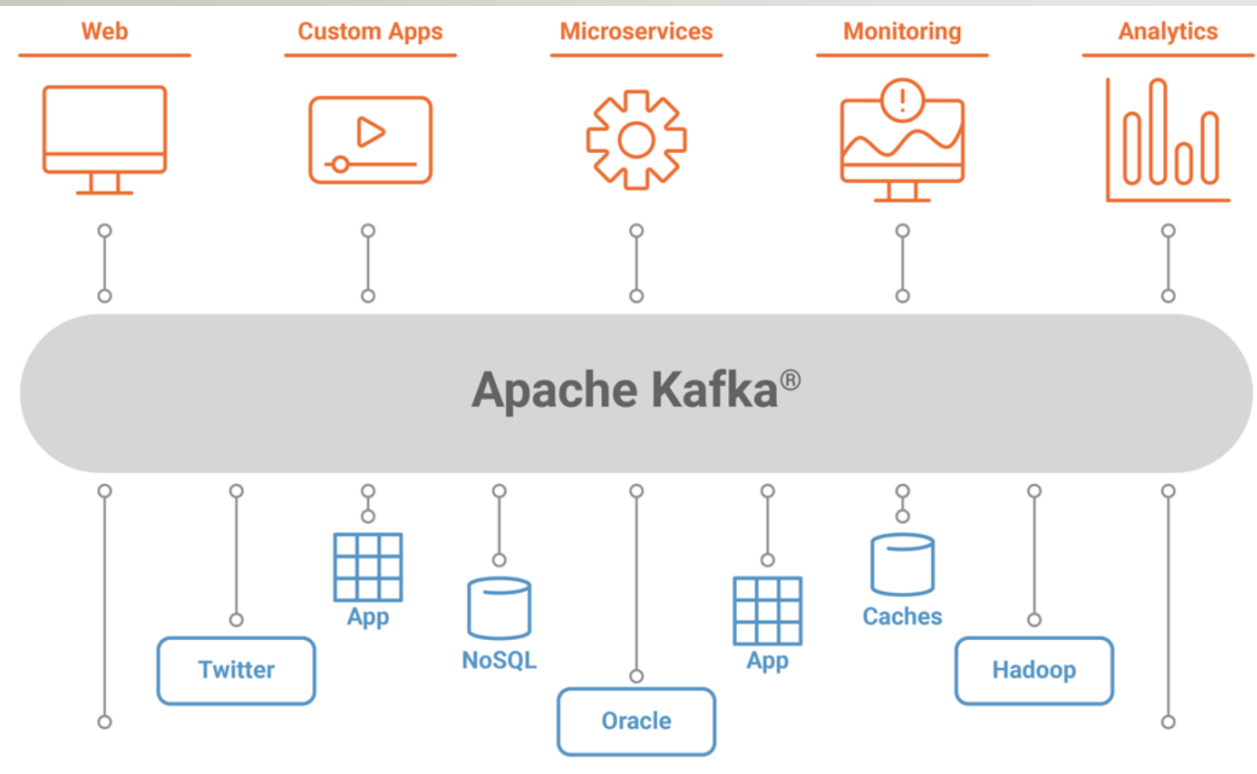
KAFKA TOPICS

- What are Kafka Partitions?
- Topics are broken down into a number of partitions. A single topic may have more than one partition, it is common to see topics with 100 partitions.
- The number of partitions of a topic is specified at the time of topic creation. Partitions are numbered starting from 0 to N-1, where N is the number of partitions. The figure below shows a topic with three partitions, with messages being appended to the end of each one



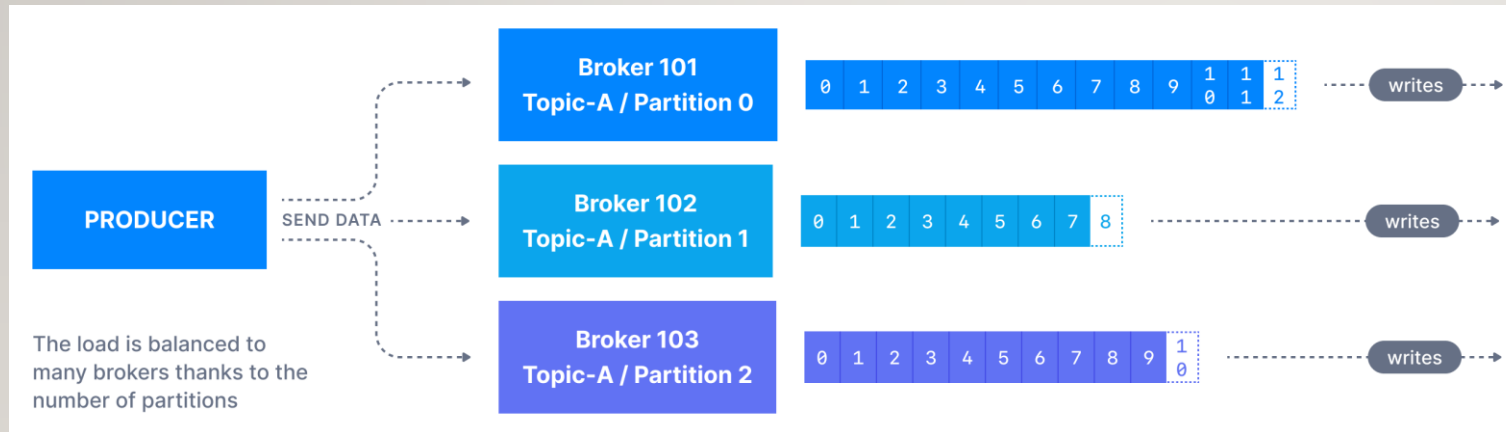
KAFKA

- Apache Kafka is an open-source distributed event streaming platform used by thousands of companies for high-performance data pipelines, streaming analytics, data integration, and mission-critical applications



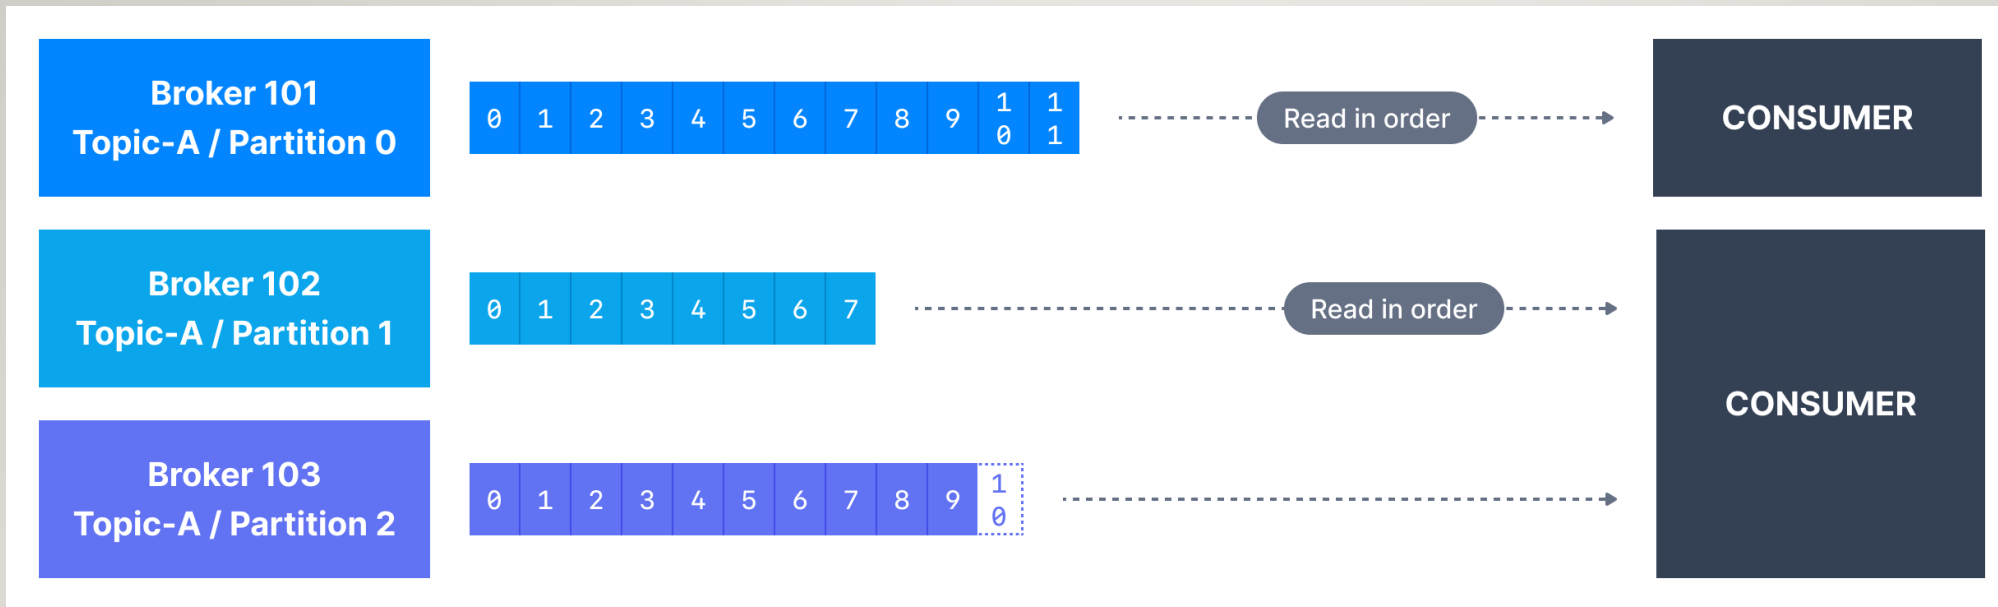
KAFKA PRODUCERS

- Applications that send data into topics are known as Kafka producers. Applications typically integrate a Kafka client library to write to Apache Kafka. Excellent client libraries exist for almost all programming languages that are popular today including Python, Java, Go, and others.
- A Kafka producer sends messages to a topic, and messages are distributed to partitions according to a mechanism such as key hashing



KAFKA CONSUMER

- Applications that read data from Kafka topics are known as consumers. Applications integrate a Kafka client library to read from Apache Kafka. Excellent client libraries exist for almost all programming languages that are popular today including Python, Java, Go, and others.
- Consumers can read from one or more partitions at a time in Apache Kafka, and data is read in order **within each partition** as shown below.



MICROSERVICES ARCHITECTURE PATTERN - SAGA



WHAT IS A DISTRIBUTED TRANSACTION?

- In the world of microservices, each service has its own local database and some operations we do in microservices may involve multiple database operations for multiple services.
- So, a distributed transaction may contain local transactions for multiple databases.
- What we often encounter is the need to guarantee data consistency, updated in parallel by different actors.



WHAT IS SAGA?

- Saga first appeared in a paper published by Hector Garcia-Molina & Kenneth Salem in 1987.
- The core idea is to split a long transaction into multiple short transactions, coordinated by the Saga transaction coordinator, with the global transaction completing normally if each short transaction completes successfully, and invoking the compensating operations one at a time according to the reverse order if a step fails.
- The Saga pattern is a widely used pattern for distributed transactions. It is asynchronous and reactive.

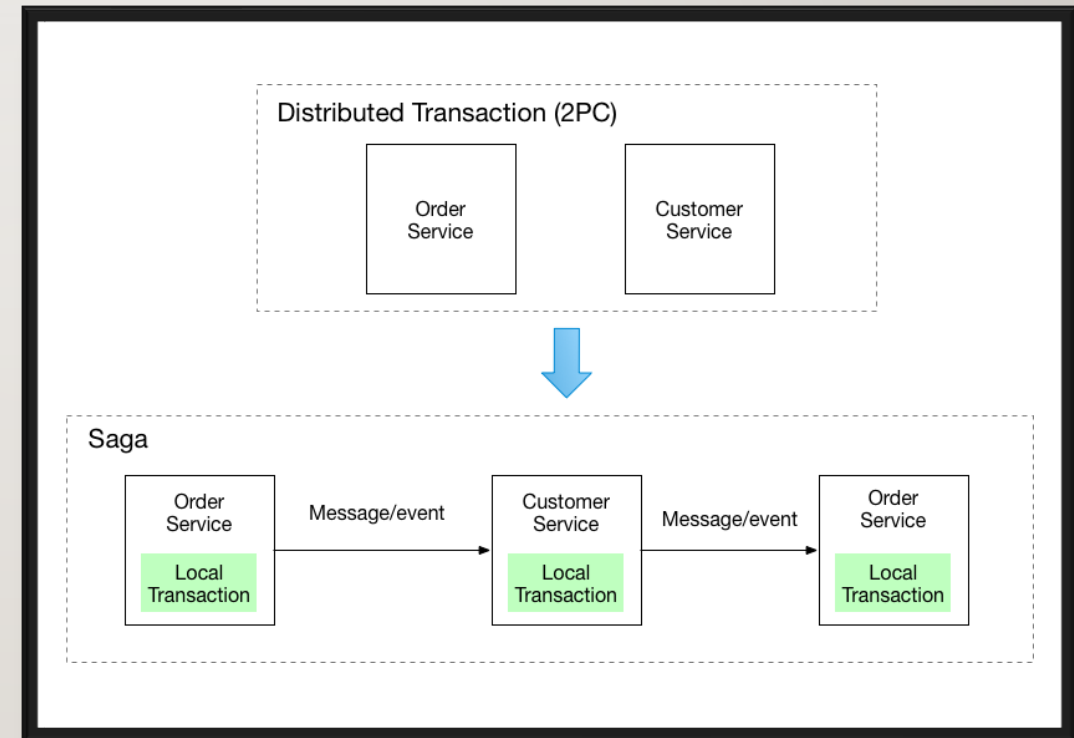


PATTERN: SAGA

- Context
 - You have applied the Database per Service pattern. Each service has its own database. Some business transactions, however, span multiple service so you need a mechanism to implement transactions that span services.
 - For example, let's imagine that you are building an e-commerce store where customers have a credit limit. The application must ensure that a new order will not exceed the customer's credit limit.
 - Since Orders and Customers are in different databases owned by different services the application cannot simply use a local ACID transaction.

PATTERN: SAGA

- Implement each business transaction that spans multiple services as a saga.
- A saga is a sequence of local transactions.
- Each local transaction updates the database and publishes a message or event to trigger the next local transaction in the saga.
- If a local transaction fails because it violates a business rule then the saga executes a series of compensating transactions that undo the changes that were made by the preceding local transactions.

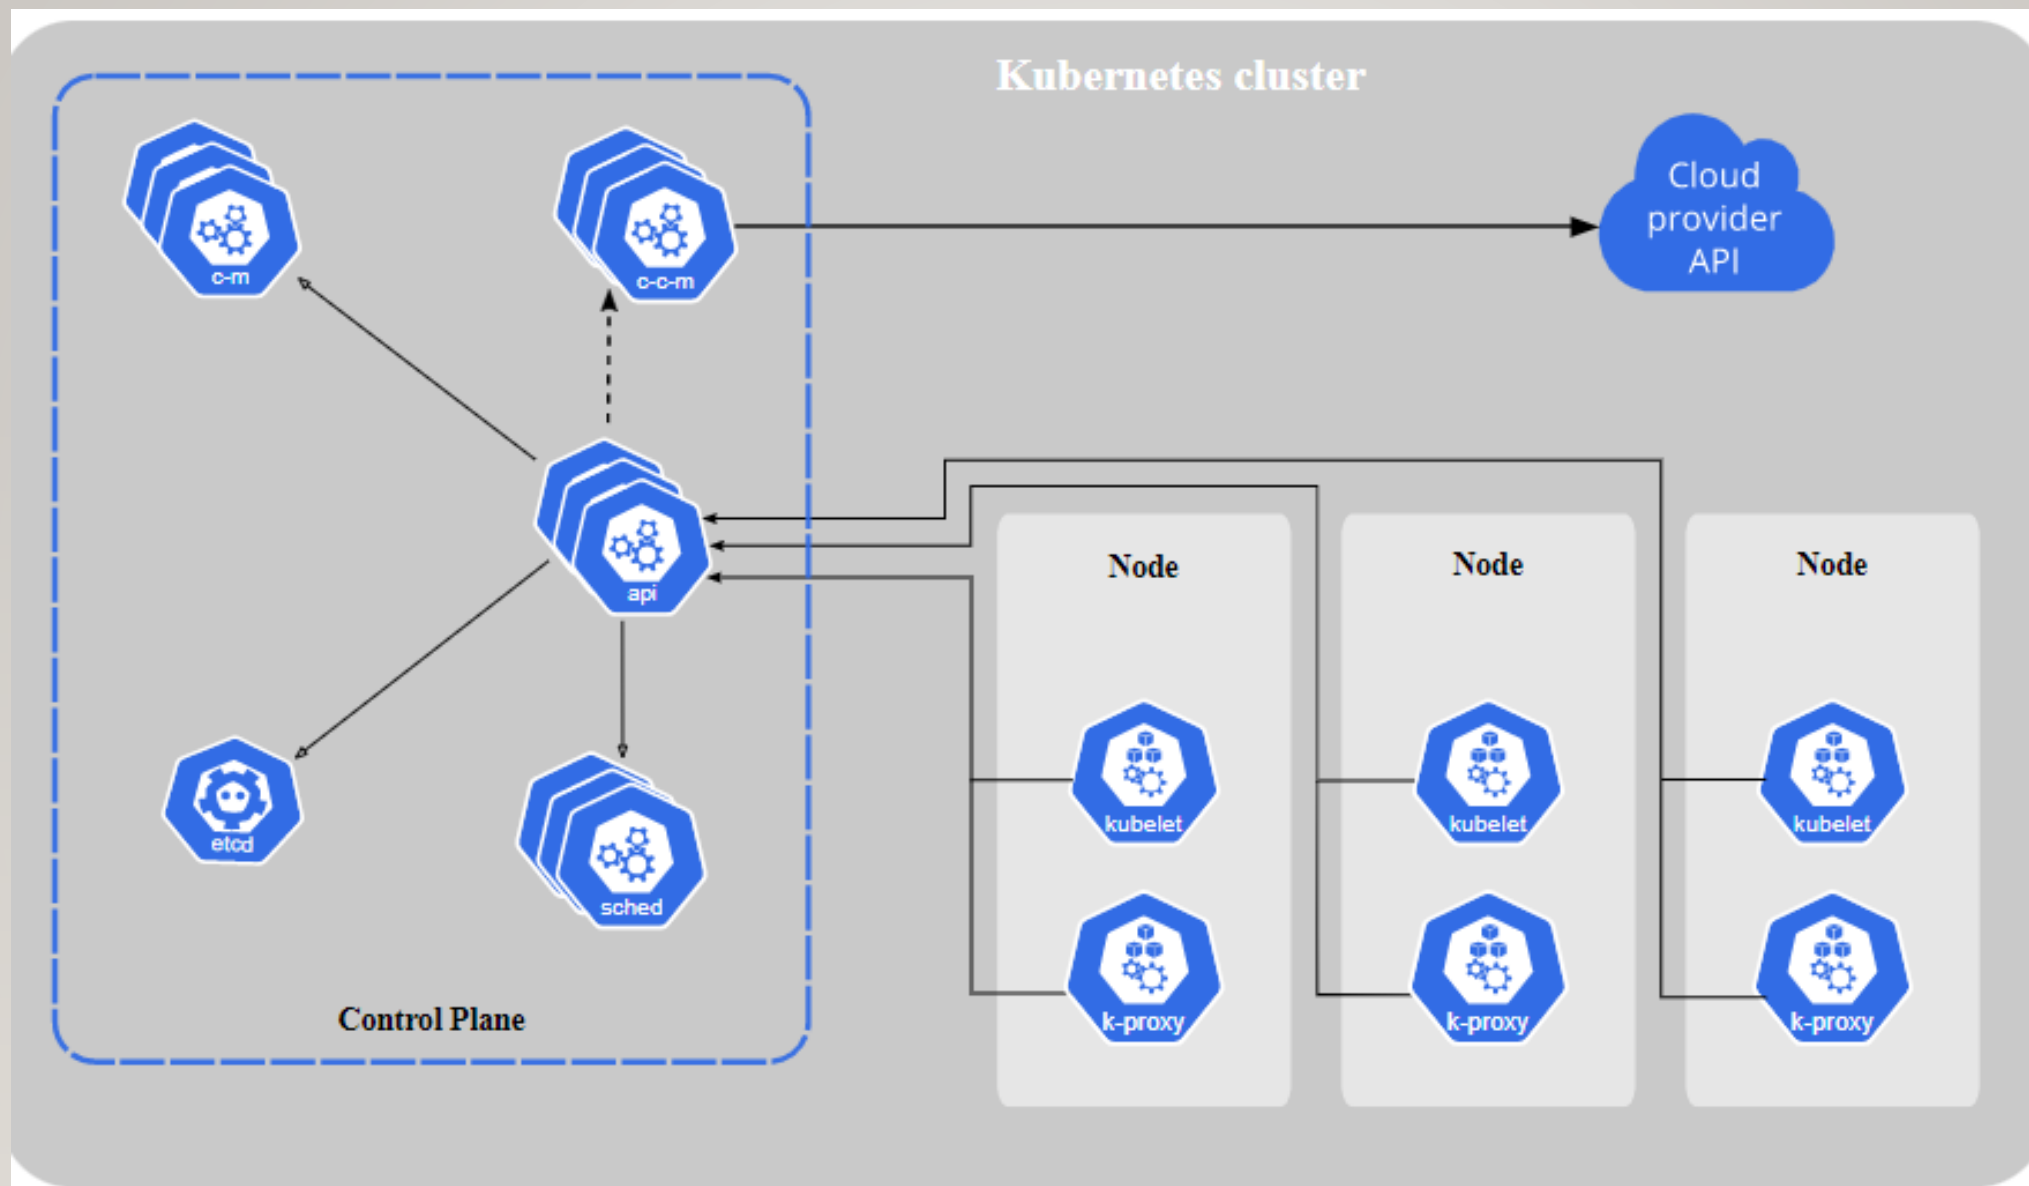


WHAT IS DTM?

- DTM is a distributed transaction framework which provides cross-service eventual data consistency.
- It provides saga, tcc, xa, 2-phase message strategies for a variety of application scenarios.
- It also supports multiple languages and multiple store engine to form up a transaction.
- There are some features of DTM:
 - Extremely easy to adopt
 - Easy to use
 - Multi-language support
 - Easy to deploy, easy to extend
 - Multiple distributed transaction protocols
- `docker run -d -p 36789:36789 -p 36790:36790 --name=dtm-svc yedf/dtm:1.13`

KUBERNETES





AKS / EKS

